

# Concurrency, Parallelism and Distribution (CPD)

Concurrency:  
Lists, Parallelism and Dynamic Code Loading in Erlang

Joaquim Gabarro  
([gabarro@cs.upc.edu](mailto:gabarro@cs.upc.edu))

Computer Science  
BarcelonaTech, Universitat Politècnica de Catalunya

# Basic reading

(1) Joe Armstrong

Programming Erlang, Software for a Concurrent World

Pragmatic Bookshelf, 2007.

Download the code from:

<https://pragprog.com/book/jaerlang/programming-erlang>

(2) Joe Armstrong, Erlang

CACM, September 2010, Vol 53, No 9, 68-75

(3) Jim Larson, Erlang for Concurrent Programming

CACM, March 2009, Vol 52, No 3, 48-56

Slides are based on Armstrong's book

Lists

Execution time & Straightforward parallelism

More on messages and processes

Dynamic code loading

# 1. Lists

# Lists

- ▶ We **use** lists to store **variable numbers** of things.
- ▶ We **create** a list by enclosing the list elements in **square brackets** and separating them with **commas**.

Example: Shopping list:

```
1> ThingsToBuy = [{apples,10},{pears,6},{milk,3}].  
[ {apples,10}, {pears,6}, {milk,3} ]
```

- ▶ The **first element** of a list is the **head** of the list.
- ▶ If you **remove the head**, what's left is the **tail** of the list.

# Defining Lists (1)

- ▶ If  $T$  is a list,  $[H|T]$  is also a list, with head  $H$  and tail  $T$ .
- ▶ The vertical bar  $|$  separates the head from its tail.
- ▶  $[]$  is the empty list.

```
7> [a| [b,c]].
```

```
[a,b,c]
```

```
8> [[2]| [b,c]].
```

```
[[2],b,c]
```

```
9> [2|b,c].
```

```
* 1: syntax error before: ','
```

```
9> [a|[]].
```

```
[a]
```

## Defining Lists (2)

We can add more than one element to the beginning of **T** by writing **[E1,E2,..,En|T]**.

```
3> ThingsToBuy1 = [{oranges,4},{newspaper,1}|ThingsToBuy].  
[ {oranges,4},{newspaper,1},{apples,10},{pears,6},{milk,3}]
```

# Extracting Elements

If we have a **nonempty** list a **L**, then the expression  $[X|Y] = L$ , where **X** and **Y** are **unbound variables**, will extract the **head** of the list into **X** and the **tail** of the list into **Y**.

Example: We have our shopping list ThingsToBuy1

```
13> [Buy1|ThingsToBuy2] = ThingsToBuy1.  
[{oranges,4},{newspaper,1},{apples,10},{pears,6},{milk,3}]  
14> Buy1.  
{oranges,4}  
15> ThingsToBuy2.  
[{newspaper,1},{apples,10},{pears,6},{milk,3}]
```



# Simple list processing (mylists.erl)

`sum(L)` computes the sum of the elements in `L`

```
sum([H|T]) -> H + sum(T);  
sum([]) -> 0.
```

**Execution example:**

```
16> cd("D:/.../Erlang_code").  
..  
17> c(mylists)  
...  
22> L=[1,2,3].  
[1,2,3]  
23> mylists:sum(L).  
6
```

## Common list processing functions (mylists.erl)

`map(F, L)` returns a list where every element is obtained applying `F` to the corresponding element in `L`.

```
map(_, []) -> [];  
map(F, [H|T]) -> [F(H)|map(F, T)].
```

`member(X, L)` returns `true` if `X` is an element of `L`, `false` otherwise.

```
member(H, [H|_]) -> true;  
member(H, [_|T]) -> member(H, T);  
member(_, []) -> false.
```

`reverse(L)` reverses the order of the elements in `L`

```
reverse(L) -> reverse(L, []).  
reverse([H|T], L) -> reverse(T, [H|L]);  
reverse([], L) -> L.
```

# Strings and Lists (1)

When the shell prints the value of a list it **prints the list as a string**, but **only if all the integers represent printable values**.

```
1> L1= [1,2,3,83,117,114,112,114,105,115,101] .  
...  
2> L1.  
[1,2,3,83,117,114,112,114,105,115,101]  
3> L2= [83,117,114,112,114,105,115,101] .  
...  
6> L2.  
"Surprise"
```

American Standard Code for Information Interchange, ASCII:

$83 \rightarrow \text{S}, 117 \rightarrow \text{u}, \dots, 101 \rightarrow \text{3}$

## Strings and Lists (2)

`sep(L, N)` returns  $\{L1, L2\}$  so that  $L1 ++ L2 == L$ ,  $\text{length}(L1) = N$ .

```
sep(L, 0) -> {[], L};  
sep([H|T], N) ->  
    {L1, L2} = sep(T, N-1),  
    {[H|L1], L2}.
```

The call `sep(L, 3)` returns  $\{[1, 2, 3], \text{"Surprise"}\}$

# List comprehension

- ▶ Is construct for creating a list based on existing lists.
- ▶ It follows the form of the mathematical set-builder notation in the Zermelo-Frankel theory.

Example: set-builder notation:

$$S = \{ 2 \cdot x \mid x \in \{0, \dots, 100\}, x^2 > 3, x^2 < 50 \}$$

$S = \{ \underbrace{2 \cdot x}_{\text{output function}} \mid \underbrace{x}_{\text{variable}} \in \underbrace{\{0, \dots, 100\}}_{\text{input set}}, \underbrace{x^2 > 3}_{\text{predicate}}, \underbrace{x^2 < 50}_{\text{predicate}} \}$

Erlang translation

```
24> S = [2*X || X <- lists:seq(0,100), X*X > 3, X*X<50].  
[4,6,8,10,12,14]
```

# List comprehension, examples

```
27> [{X,Y} || X<-[1,2,3], Y<-[a,b,c]].  
[{1,a},{1,b},{1,c},{2,a},{2,b},{2,c},{3,a},{3,b},{3,c}]  
28> [X || X<-[1,2,3,4], Y<-[3,4,5,6], X==Y].  
[3,4]  
29> [X+1 || X<-[1,2,3,4], X rem 2==0].  
[3,5]
```

**General form** [Exp || Gen1, ..., GenN, Filter1, ..., FilterN]

**Example: permutations**

```
-module(perm).  
-compile([export_all]).  
  
perm([])->[[]];  
perm(L)->[[H|T] || H<-L, T<-perm(L--[H])].
```

**Execution example:**

```
45> perm:perm([c,a,t]).  
[[c,a,t],[c,t,a],[a,c,t],[a,t,c],[t,c,a],[t,a,c]]
```

# Sequential quicksort

```
-module(pqs).  
-compile([export_all]).  
  
%% sequential quicksort  
qs([]) -> [];  
qs([H|T]) ->  
    LT = [X || X <- T, X < H],  
    GE = [X || X <- T, X >= H],  
    qs(LT) ++ [H] ++ qs(GE).
```

**Example: Given [27,82,43,15,10,38,9,3]**

```
[H|T]=[27|82,43,15,10,38,9,3]  
LT=[15,10,9,3], GE=[82,43,38]  
qs(LT)=[3,9,10,15], qs(GE)=[38,43,82]  
qs(LT)++[H]++qs(GE)=[3,9,10,15,27,38,43,82]
```

## 2. Execution time & Straightforward parallelism



# Execution time

`apply(M, F, [Arg1,...,ArgN])` equivalent `M:F(Arg1,...,ArgN)`

`chrono(M, F, P)` returns the computing time of `M:F` having `P` as parameters.

```
-module(new_pqs).  
-compile([export_all]).  
....  
chrono(M, F, P) ->  
    {_, Seconds, Micros} = erlang:timestamp(),  
    T1 = Seconds + (Micros/1000000.0),  
    apply(M, F, P),  
    {_, Seconds2, Micros2} = erlang:timestamp(),  
    T2 = Seconds2 + (Micros2/1000000.0),  
    T2 - T1.
```

## qs execution time

Macro `?MODULE` expands to the current module name.

```
-module(pqs).  
-compile([export_all]).  
  
%% test functions  
random_list(N) -> random_list(N, N, []).  
random_list(0, _, L) -> L;  
random_list(N, M, L)  
    -> random_list(N-1, M, [random:uniform(M) | L]).  
  
test_seq(N) -> L=random_list(N), chrono(?MODULE, qs, [L]).
```

### Execute a test

```
2> c(new_pqs).  
{ok,new_pqs}  
3> new_pqs:test_seq(5000000).  
9.9530000000037719
```

# Parallel quicksort pqs

```
...  
%% parallel quicksort  
pqs(L) -> P = spawn(?MODULE, pqs2, [self(), L]), rcv(P).  
  
rcv(P) ->receive {P, X} -> X end.  
  
pqs2(P, L) ->  
    if  
        length(L) < 100000 ->  
            P ! {self(), qs(L)};  
        true ->  
            [H | T] = L,  
            LT = [X || X <- T, X < H],  
            GE = [X || X <- T, X >= H],  
            P1 = spawn(?MODULE, pqs2, [self(), LT]),  
            P2 = spawn(?MODULE, pqs2, [self(), GE]),  
            L1 = rcv(P1),  
            L2 = rcv(P2),  
            P ! {self(), L1 ++ [H] ++ L2}  
    end.
```

## qs versus pqs

Explicit parallelism does not help very much:

```
2> c(new_pqs) .  
{ok,new_pqs}  
3> new_pqs:test_seq(5000000) .  
9.9530000000037719  
4> new_pqs:test_par(5000000) .  
8.1720000000020489
```

## 4. More on messages and processes

# Receive with a Timeout

A **receive** statement might **wait forever** for a message that never comes.

```
receive
  Pattern1 [when Guard1] -> Expressions1;
  Pattern2 [when Guard2] ->Expressions2;
  ...
after Time ->
  Expressions
end.
```

If **no matching message** has arrived within **Time milliseconds** of entering the receive expression, then the process will **stop waiting for a message** and evaluate **Expressions**.

## Example : Sleep

Function `sleep(T)` **suspends** the current process for `T` milliseconds.

Download `lib_misc.erl`

```
sleep(T) ->  
    receive  
    after T ->  
        true  
    end.
```

## Example: Receive with Timeout Value of Zero (1)

The function `flush_buffer` entirely empties the mailbox of a process.

Download `lib_misc.erl`

```
flush_buffer() ->  
    receive  
        _Any -> flush_buffer()  
    after 0 ->  
        true  
    end.
```

Without the `timeout` clause, `flush_buffer` would suspend forever and not return when the mailbox was empty.



## Example: Receive with Timeout Value of Zero (2)

We use a zero **timeout** to implement a form of **priority receive** as follows (download `lib_misc.erl`)

```
priority_receive() ->  
    receive  
        {alarm, X} -> {alarm, X}  
    after 0 ->  
        receive  
            Any -> Any  
        end  
    end.
```

- ▶ If there is **no message at all**, it will suspend in the innermost receive and return the first message it receives.
- ▶ If there is a **message matching** alarm, X, then this message will be returned immediately.
- ▶ Remember that the after section is checked only after pattern matching has been performed on all the entries.

## Example: Implementing a Timer (1)

- ▶ The function `stimer:start(Time, Fun)` will **evaluates Fun after Time** ms.
- ▶ It **returns a handle** (which is a PID) used to **cancel** the timer if required.

(download `stimer.erl`)

```
-module(stimer).  
-export([start/2, cancel/1]).
```

```
start(Time, Fun)->spawn(fun()->timer(Time, Fun) end).  
cancel(Pid) -> Pid ! cancel.
```

```
timer(Time, Fun) ->  
    receive  
        cancel ->void  
    after Time ->  
        Fun()  
    end.
```

## Example: Implementing a Timer (2)

We can test this as follows:

I wait more than five seconds so that the timer would **trigger**.

```
1> Pid = stimer:start(5000, fun() ->
                                io:format("timer event~n") end).
<0.42.0>
timer event
```

I'll start a timer and **cancel** it before has expired:

```
2> Pid1 = stimer:start(25000, fun()
                        -> io:format("timer event~n") end).
<0.49.0>
3> stimer:cancel(Pid1).
cancel
```

# Selective Receive (1)

- ▶ Each process in Erlang has an associated **mailbox**.
- ▶ When you **send** a message to the process, the message is **put** into the **mailbox**.
- ▶ The only time the mailbox is **examined** is when your program evaluates a **receive**. statement.

```
receive
    Pattern1 [when Guard1] -> Expressions1;
    Pattern2 [when Guard1] -> Expressions1;
...
after
    Time -> ExpressionTimeout
end
```

## Selective Receive (2)

Receive works as follows:

- ▶ When we enter a receive statement, we start a timer.
- ▶ Take the **first message** in the mailbox and try to **match** it against **Pattern1**, **Pattern2**, and so on. If the **match** succeeds, the message is **removed** from the mailbox, and the expressions following the pattern are **evaluated**.
- ▶ If **none** of the patterns in the receive statement matches the **first** message in the mailbox, then the first message is removed from the mailbox and put into a **save queue**. The **second** message in the mailbox is then tried. This procedure is **repeated** until a matching message is found or until all the messages in the mailbox have been examined.

## Selective Receive (3)

- ▶ If **none of the messages** in the mailbox matches, then the process is suspended and will be rescheduled for execution the next time a new message is put in the mailbox. Note that when a new message arrives, the messages in the save queue are not rematched.
- ▶ As soon as a **message has been matched**, then all messages that have been put into the save queue are reentered into the mailbox in the order in which they arrived at the process. If a timer was set, it is cleared.
- ▶ If the **timer elapses** when we are waiting for a message, then evaluate the expressions `ExpressionsTimeout` and put any saved messages back into the mailbox in the order in which they arrived at the process.

# Registered Processes

Erlang has a method for **publishing** a process identifier so that any process in the system can communicate with this process.

`register(AnAtom, Pid)`. Register the process `Pid` with the name `AnAtom`. The registration fails if `AnAtom` has already been used to register a process.

```
1> Pid = spawn(fun area_server0:loop/0).
<0.51.0>
2> register(area, Pid).
true
3> area ! {rectangle, 4, 5}.
Area of rectangle is 20
{rectangle,4,5}
```

# A Clock (1)

We use register to make a registered process that represents a clock (download `clock.erl`).

```
-module(clock).  
-export([start/2, stop/0]).  
  
start(Time, Fun) ->  
    register(clock, spawn(fun() -> tick(Time, Fun) end)).  
  
stop() -> clock ! stop.  
  
tick(Time, Fun) ->  
    receive  
        stop -> void  
    after Time ->  
        Fun(), tick(Time, Fun)  
    end.
```



## A Clock (2)

The clock will happily tick away until you stop it:

```
3> clock:start(5000, fun()  
    -> io:format("TICK ~p~n",[erlang:now()]) end).  
true  
TICK {1350,553188,828000}  
TICK {1350,553193,843000}  
TICK {1350,553198,859000}  
TICK {1350,553203,874000}  
4> clock:stop().  
stop
```

## 5. Dynamic code loading

# Dynamic code loading: simple idea

Dynamic code loading is one surprising features of Erlang.

The idea is simple: every time we call

```
someModule:someFunction(...)
```

we'll always call the latest version of the function in the latest version of the module, even if we recompile the module while code is running in this module.

## Example: a and b (1)

Consider two little modules **a** and **b**.

- ▶ If **a** calls **b** in a loop and **we recompile b**, then **a** will automatically **call the new version** of **b** the next time **b** is called.
- ▶ If **different processes a** are running and all of them call **b**, then **all of them** will call the **new version** of **b** if **b** is recompiled.

Module **b** is very simple

```
-module(b) .  
-export([x/0]) .  
  
x() -> 1.
```

## Example: Module a (2)

The `a` processes sleep for three seconds, wake up and call `b:x()`, and then print the result.

```
-module(a) .  
-compile(export_all) .  
  
start(Tag) -> spawn(fun() -> loop(Tag) end) .  
  
loop(Tag) ->  
    sleep(),  
    Val = b:x(),  
    io:format("Vsn1 (~p) b:x() = ~p~n", [Tag, Val]),  
    loop(Tag) .  
  
sleep() ->  
    receive  
        after 15000 -> true  
    end.
```

## Example: Start up scenario (3)

**Start up scenario:** Compile **b** and **a** and start a couple of **a** processes:

```
2> c(b) .  
{ok,b}  
3> c(a) .  
{ok,a}  
4> a:start(one) .  
<0.53.0>  
Vsn1 (one) b:x() = 1  
5> a:start(two) .  
<0.55.0>  
Vsn1 (one) b:x() = 1  
Vsn1 (two) b:x() = 1  
Vsn1 (one) b:x() = 1  
Vsn1 (two) b:x() = 1  
Vsn1 (one) b:x() = 1  
Vsn1 (two) b:x() = 1
```

## Example: Change and recompile b (4)

**Scenario:** Go into the **editor** and change **b** to the following:

```
-module(b) .  
-export([x/0]) .  
x() -> 2.
```

We **recompile** **b** in the shell. This is what happens:

```
6> c(b) .  
{ok,b}  
Vsn1 (two) b:x() = 2  
Vsn1 (one) b:x() = 2  
Vsn1 (two) b:x() = 2  
...
```

The **two original versions** of **a** are still running, but they call the **new version** of **b**. When we call **b:x()** from within the module **a**, we really call **the latest version of b**.

## Example: Change and recompile a (5)

**Scenario:** Go into the **editor** and change the version **a** from

```
io:format("Vsn1 (~p) b:x() = ~p~n",[Tag, Val]),
```

into

```
io:format("Vsn2 (~p) b:x() = ~p~n",[Tag, Val]),
```

Now we compile and start a new **a** and the result is:



## Example: ... and the result is (6)

```
7> c(a) .
{ok,a}
Vsn1 (one) b:x() = 2
Vsn1 (two) b:x() = 2
8> a:start(three) .
<0.67.0>
Vsn1 (one) b:x() = 2
Vsn1 (two) b:x() = 2
Vsn2 (three) b:x() = 2
Vsn1 (one) b:x() = 2
Vsn1 (two) b:x() = 2
Vsn2 (three) b:x() = 2
...
```

When we start the **new version** of **a**, we see that **new version running**. The existing processes running the **first version** of **a** **are still running** that old version.

## Example: Change and recompile b again (7)

**Scenario:** Go into the **editor** and change **b** to the following:

```
-module(b) .  
-export([x/0]) .  
x() -> 3.
```

We **recompile b** in the shell. This is what happens:

```
9> c(b) .  
{ok,b}  
Vsn1 (one) b:x() = 3  
Vsn1 (two) b:x() = 3  
Vsn2 (three) b:x() = 3  
....
```

Both the **old and new** versions of **a** call the **latest** version of **b**.

## Example: Change and recompile a Vsn3 (7)

**Scenario:** Recompile **a**. into a new Vsn3 and start a new **a**.

```
10> c(a).  
{ok,a}  
11> a:start(four).  
<0.79.0>  
Vsn2 (three) b:x() = 3  
Vsn3 (four) b:x() = 3  
Vsn2 (three) b:x() = 3  
Vsn3 (four) b:x() = 3  
...
```

The output contains the **last two** versions of **a** (versions 2 and 3); the process running **version 1** of a's code has **died**.

Erlang can have two versions of a module running at any one time, the current version and an old version.

That's all!