

# CUDA BASICS

# CUDA Basics

## ■ Outline

- GPU architecture
- Devices, kernel definition and offloading
- Blocks, threads and indexing
- Accessing (global) memory
- Cooperating threads and shared memory
- Reductions in CUDA

# CPU vs. GPU architecture

## ■ CPUs are designed for general-purpose computing

- Sophisticated control to exploit ILP. ALU to exploit DLP.
- Multicores (independent control flow) for TLP
- Large caches to reduce impact of long latency memory accesses and to enable sharing in multicores



# CPU vs. GPU architecture

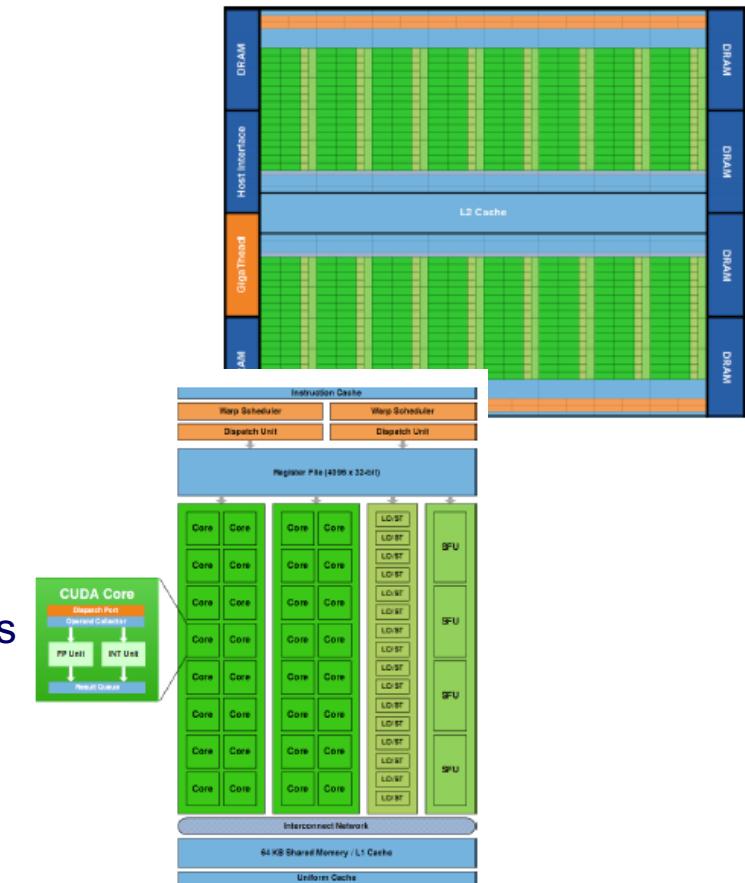
- GPUs are specialized for highly parallel, compute-intensive computation
  - Simple control: same computation on all compute units
  - Massive number of threads to reduce impact of long latency memory accesses



# NVIDIA GPU architecture: Parameters, terminology ....

## ■ NVIDIA Fermi

- 16 streaming multiprocessors (SM) interconnected via a 768k L2 cache crossbar
- Each SM has 32 cores and 4096 registers, sharing 64 KB (48K L1 and 16K shared memory or viceversa)
- Each core has one FP unit and one integer unit
- Each SM has 16 load/store units, allowing sixteen threads per clock to calculate memory addresses
- Each SM has 4 Special Function Units (SFU) that execute complex instructions (sin, cosine, reciprocal, and square root)



# NVIDIA GPU architectures

## ■ Exercise:

What are other GPU architectures ?

Identify the architectural parameters for

NVIDIA Kepler

NVIDIA Pascal

NVIDIA Volta

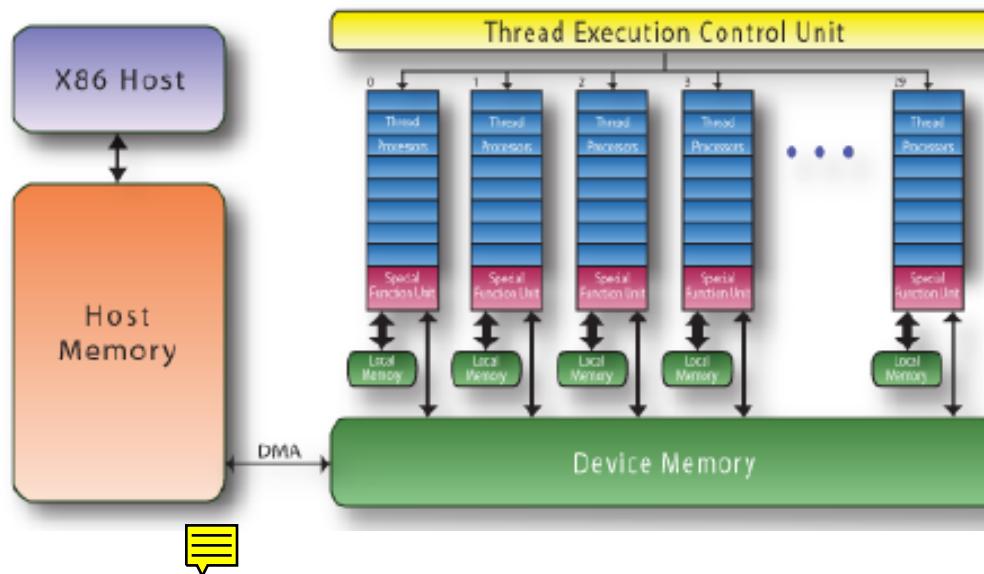
## ■ Exercise:

What is the GPU architecture for the LAB machines ?

What are its main architectural parameters ?

# CPU-GPU block diagram

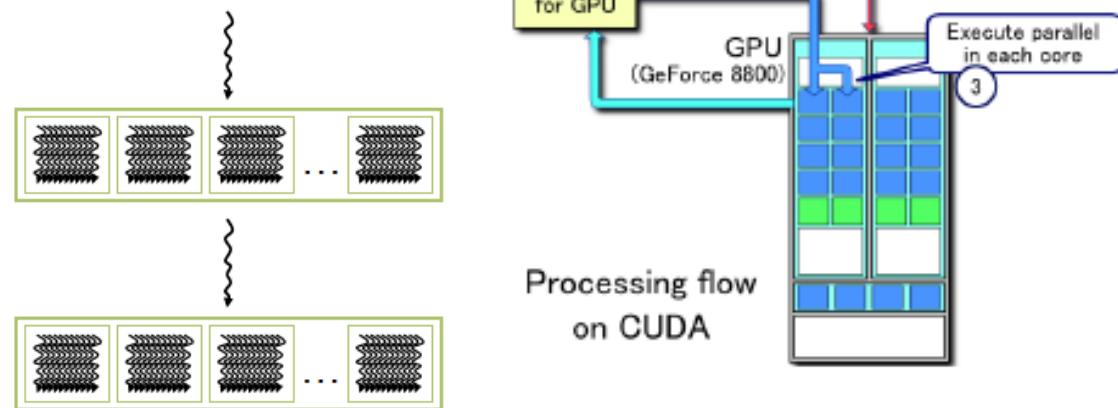
- Separate address spaces between CPU (host memory) and GPU (device memory).  
Communication using IO commands and DMA memory transfers (e.g. PCI Express)



# CPU-GPU execution flow

## ■ CPUs for sequential or modestly parallel parts, GPUs for highly parallel parts (kernels in CUDA)

- 1) Copy processing data
- 2) Instruct the processing
- 3) Execute parallel computation
- 4) Copy back the result



# CUDA Basics

## ■ Outline

- GPU architecture
- Devices, kernel definition and offloading
- Blocks, threads and indexing
- Accessing (global) memory
- Cooperating threads and shared memory
- Reductions in CUDA

# Device management

## ■ Application can query and select GPUs

```
cudaGetDeviceCount(int *count)  
cudaSetDevice(int device)  
cudaGetDeviceProperties(cudaDeviceProp *prop, int device)  
cudaGetDevice(int *device)
```

## ■ Device sharing

- Multiple threads can share a device
- A single thread can manage multiple devices

# Offloading kernel execution

## ■ Kernel

- Function that runs on the device and is called from host/device

## ■ Definition

- Using keyword `__global__`

```
__global__ void helloworld(void) {  
}
```

# Offloading kernel execution (cont.)

## ■ Kernel offloading for execution on GPU device

...

```
helloworld<<<1,1>>>();
```

...

## ■ Triple angle brackets mark a "kernel launch" (call from host code to device code)

- We will cover the parameters inside the brackets in a few slides
- The kernel launch is asynchronous (i.e. the host does not wait for the termination of the kernel execution, control returns to the CPU immediately)

# Offloading kernel execution (cont.)

## ■ Host may need to wait before continuing execution

- The execution of `cudaDeviceSynchronize()` blocks the CPU until all preceding CUDA kernels have completed

```
__global__ void helloworld(void) {  
}  
  
int main(void) {  
  
    helloworld<<<1,1>>>();  
    cudaThreadSynchronize();  
  
    printf("Hello World!\n");  
    return 0;  
}
```

# Offloading to multi-GPU nodes

- Any host thread can access all GPUs in the node (if more than one available) using the `cudaSetDevice` call

```
__global__ void helloworld(void) {  
  
}  
int main(void) {  
    int numDevs;  
    cudaGetDeviceCount(&numDevs);  
    for (int d = 0; d < numDevs; d++) {  
        cudaSetDevice(d);  
        helloWorld<<<1,1>>>();  
    }  
    for (int d = 0; d < numDevs; d++) {  
        cudaSetDevice(d);  
        cudaThreadSynchronize();  
    }  
    return 0;  
}
```

# CUDA Basics

## ■ Outline

- GPU architecture
- Devices, kernel definition and offloading
- Blocks, threads and indexing
- Accessing (global) memory
- Cooperating threads and shared memory
- Reductions in CUDA

# Parallel kernel

## ■ How to make use of the multiple streaming multiprocessors (SM) and threads inside each SM?

- Kernel replication using the arguments in the kernel launch
  - ✓ Replication in different SMs

```
helloWorld<<<N,1>>>();
```

- ✓ Replication in different threads but in a single SM

```
helloWorld<<<1,N>>>();
```

## Parallel kernel (II)

- Each kernel instance can be uniquely identified, which can be used to make kernel instances cooperate in the resolution of a common problem

```
__global__ void helloWorld(char* str) {
    int idx = blockIdx.x;
    str[idx] += idx;
}

int main (int argc, char *argv[]) {
    ...
    char str[] = "Hello World!";
    for(i = 0; i < 12; i++)
        str[i] -= i;
    ...
    helloWorld<<<12, 1>>>(d_str);
    ...
}
```

# Parallel kernel (III)

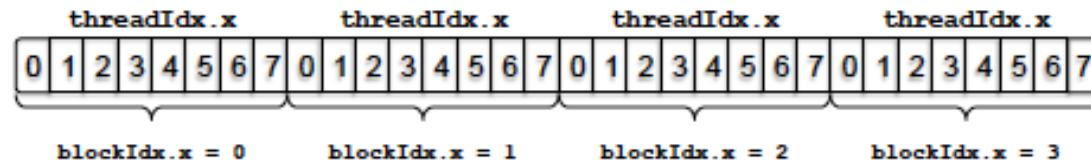
## ■ Similarly if replication occurs at the threads level

```
__global__ void helloWorld(char* str) {
    int idx = threadIdx.x;
    str[idx] += idx;
}
int main (int argc, char *argv[]) {
    ...
    char str[] = "Hello World!";
    for(i = 0; i < 12; i++)
        str[i] -= i;
    ...
    helloWorld<<<1, 12>>>(d_str);
    ...
}
```

# Thread and Data Association

## ■ Indexing a vector with blocks and threads

- Consider indexing an array with one element per thread and M threads per block (e.g. 8)



- Then a unique index for each thread is given by

```
int index = threadIdx.x + blockIdx.x * M;
```

# Example: vector addition (sequential code)

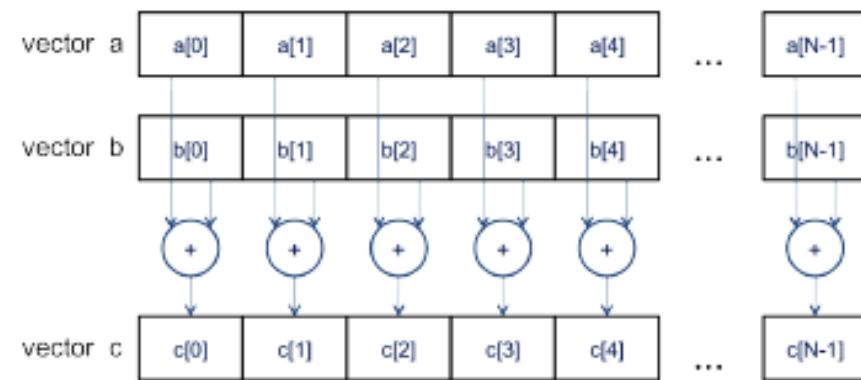
```
// Compute vector sum C = A+B
void vecAdd(float* a, float* b, float* c, int n) f
    for (i = 0, i < n, i++)
        c[i] = a[i] + b[i];
}

int main() f
    float a[N], b[N], c[N];

    // initialize a and b

    vecAdd(a, b, c, N);

    // display the results
    return 0;
}
```



# Example: vector addition (kernel definition and invocation)

## ■ Kernel code on device

```
// Each kernel invocation performs one pair-wise addition
__global__ void vecAddKernel(float* a_d, float* b_d, float* c_d, int n){

    int i = threadIdx.x + blockDim.x * blockIdx.x;
    c_d[i] = a_d[i] + b_d[i];

}
```

## ■ Kernel invocation on host

```
__host__ int vectAdd(float* a, float* b, float* c, int n) {
    ...
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<<ceil(n/256), 256>>>(a, b, c, n);
    ...
}
```

# Function declarations in CUDA

- Functions in CUDA programs can be:

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host <b>AND DEVICE !!!</b>
<code>__host__ float HostFunc()</code>	host	host

- A kernel function must return `void`

# Example: vector addition (kernel code revisited)

## ■ What if blockDim.x does not divide N?

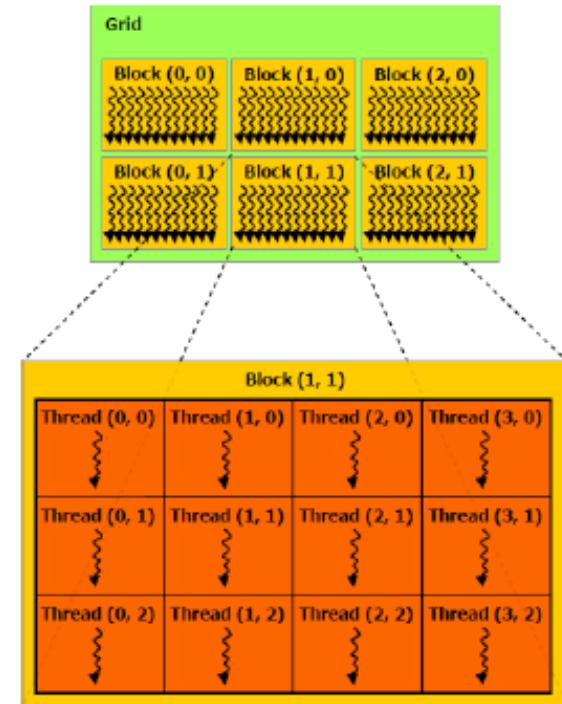
```
// Each thread performs one pair-wise addition
__global__ void vecAddKernel(float* a_d, float* b_d, float* c_d, int n) {

    int i = threadIdx.x + blockDim.x * blockIdx.x;
    if (i < n)
        c_d[i] = a_d[i] + b_d[i];
}
```

## ■ Guarded execution to ensure that we are not going out from the N elements

# Identifying threads inside a grid

- In general, the GPU offers a grid of threads, divided into blocks (thread blocks), and each block is further divided into threads
- The grid of thread blocks can actually be partitioned into 1, 2 or 3 dimensions



# Identifying threads inside a grid

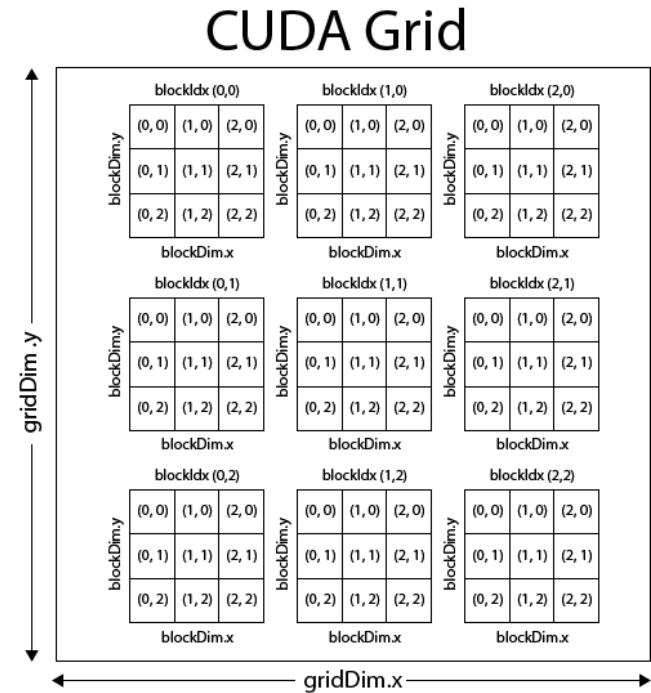
- Each thread uses indices (1D, 2D or 3D) to decide what to do and data to work on (data decomposition)

blockIdx

threadIdx

- defined inside:

- a grid with `gridDim` blocks,
- each block with `blockDim` threads



# dim3 datatype

## ■ Data type to define variables to hold block and grid dimensionalities

```
__host__ int vectAdd(float* a, float* b, float* c, int n) {  
    ...  
    // Run ceil(n/256) blocks of 256 threads each  
    dim3 DimGrid(ceil(n/256), 1, 1);  
    dim3 DimBlock(256, 1, 1);  
    vecAddKernel<<<DimGrid,DimBlock>>>(a, b, c, n);  
    ...  
}
```

# Thread scheduling and execution

## ■ Threads in a block are divided in 32-thread warps (**scheduling units in SM**)

- Example:

if 3 blocks assigned to an SM, each with 256 threads, how many warps are there in an SM?

- Each block is divided into  $256/32 = 8$  warps
- There are  $8 \times 3 = 24$  warps

✓ At any point in time, only one of the 24 warps will be selected for instruction fetch and execution (multithreading)

# CUDA Basics

## ■ Outline

- GPU architecture
- Devices, kernel definition and offloading
- Blocks, threads and indexing
- Accessing (global) memory
- Cooperating threads and shared memory
- Reductions. in CUDA

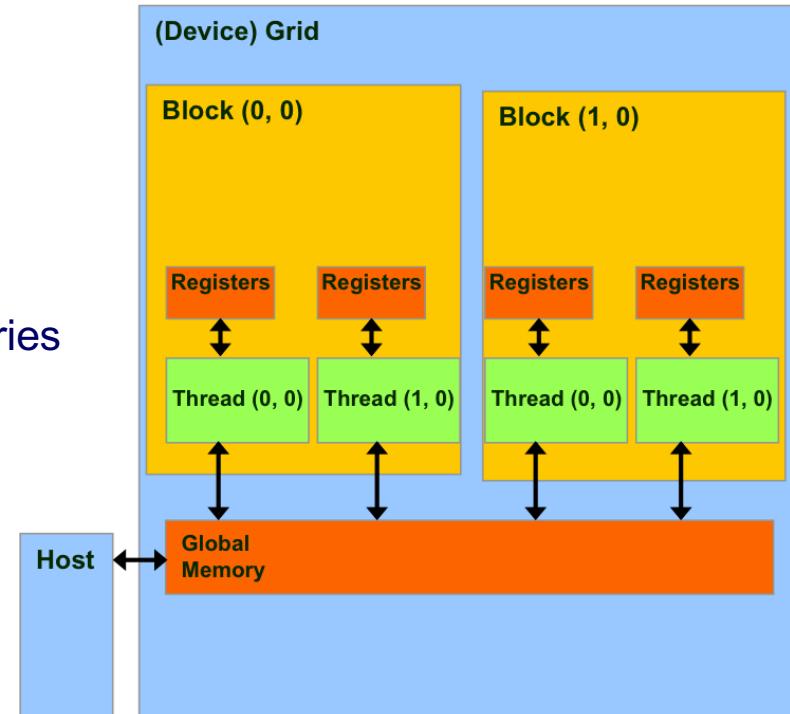
# Partial view of the device memory

## ■ Device code can

- read and write thread registers
- read and write grid global memory

## ■ Host code can

- Allocate data in grid global memory
- Transfer data between host and grid global memories



# Basic device memory management

## ■ Global memory

- Contents visible to all threads
- Variables in grid global memory declared using device attribute
- Also dynamically allocatable using `cudaMalloc` and `cudaFree`

```
#define N 1000

__device__ int A[N]; // declared outside function and kernel bodies.

// Lifetime: application
__global__ kernel() {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    A[tid]++;
}
```

# Basic device memory management

- `cudaMalloc`
  - Allocates object in the device global memory
  - Two parameters: pointer to the allocated object, size of allocated object in bytes
- `cudaFree`
  - Frees object from device global memory
  - Parameter: pointer to object

## Example: vector addition (host code)

```
int main() {
    float a[N], b[N], c[N]; // vectors in host memory
    float *dev_a, *dev_b, *dev_c; // pointers to vectors dynamically allocated in global memory

    // Allocate a, b and c on the device
    cudaMalloc( &dev_a, N * sizeof(float) ) ;
    cudaMalloc( &dev_b, N * sizeof(float) ) ;
    cudaMalloc( &dev_c, N * sizeof(float) ) ;
    ...
    // kernel invocation
    dim3 DimGrid(ceil(N/256), 1, 1);
    dim3 DimBlock(256, 1, 1);
    vecAddKernel<<<DimGrid,DimBlock>>>(dev_a, dev_b, dev_c, N);
    ...
    // Free the memory allocated on device
    cudaFree( dev_a );
    cudaFree( dev_b );
    cudaFree( dev_c );
    return 0;
}
```

# Basic device memory management (cont.)

## ■ Memory data transfer

`cudaMemcpy( . . . )`

- ✓ Synchronous
  - Blocks the CPU until the copy is complete
  - Copy begins when all preceding CUDA calls have completed
- ✓ 4 parameters: pointer to destination, pointer to source, number of bytes to be copied and type of transfer

`cudaMemcpyAsync( . . . )`

- ✓ Asynchronous, does not block the CPU, allowing the overlap of data transfer and computation
- ✓ Additional stream parameter that if not 0 can be used to synchronize

## Example: vector addition (host code) (cont'd)

```
int main() {
    float a[N], b[N], c[N]; // vectors in host memory
    float *dev_a, *dev_b, *dev_c; // pointers to vectors dynamically allocated in global memory

    cudaMalloc( &dev_a, N * sizeof(float) );
    cudaMalloc( &dev_b, N * sizeof(float) );
    cudaMalloc( &dev_c, N * sizeof(float) );

    // Initialize a and b in host memory (same as sequential)
    // Copy a and b from host to device memory
    cudaMemcpy( dev_a, a, N * sizeof(float), cudaMemcpyHostToDevice );
    cudaMemcpy( dev_b, b, N * sizeof(float), cudaMemcpyHostToDevice );

    dim3 DimGrid(ceil(N/256), 1, 1);
    dim3 DimBlock(256, 1, 1);
    vecAddKernel<<<DimGrid,DimBlock>>>(dev_a, dev_b, dev_c, N);

    // Copy back c from device to host
    cudaMemcpy( c, dev_c, N * sizeof(float), cudaMemcpyDeviceToHost );

    // Display the results

    cudaFree( dev_a ); cudaFree( dev_b ); cudaFree( dev_c );
    return 0;
}
```

**ALLOC**

**COPY**

**COMPUTE**

**COPY**

**FREE**

# Sharing data between GPUs

## ■ Options:

- Explicit copies via host
- Peer-to-peer memory access:
  - ✓ Direct copy from pointer on device to pointer on another device

```
cudaMemcpyPeer( void *dst, int dstDevice, void *src, int  
                srcDevice, size_t count )
```

- Unified Memory :
  - ✓ shared address space between devices and host (not covered in this course)

# CUDA Basics

## ■ Outline

- GPU architecture
- Devices, kernel definition and offloading
- Blocks, threads and indexing
- Accessing (global) memory
- Cooperating threads and shared memory
- Reductions in CUDA

# Synchronization: Threads vs. blocks

- In CUDA it is not possible to synchronize threads belonging to different blocks
- Unlike blocks, threads within a block have mechanisms to:
  - Synchronize: the instruction `_syncthreads()` forces all warps (i.e. all threads in a block) to wait until the rest have reached the same point
  - Communicate via memory:
    - ✓ let's take a look at a more complete view of the device memory ...

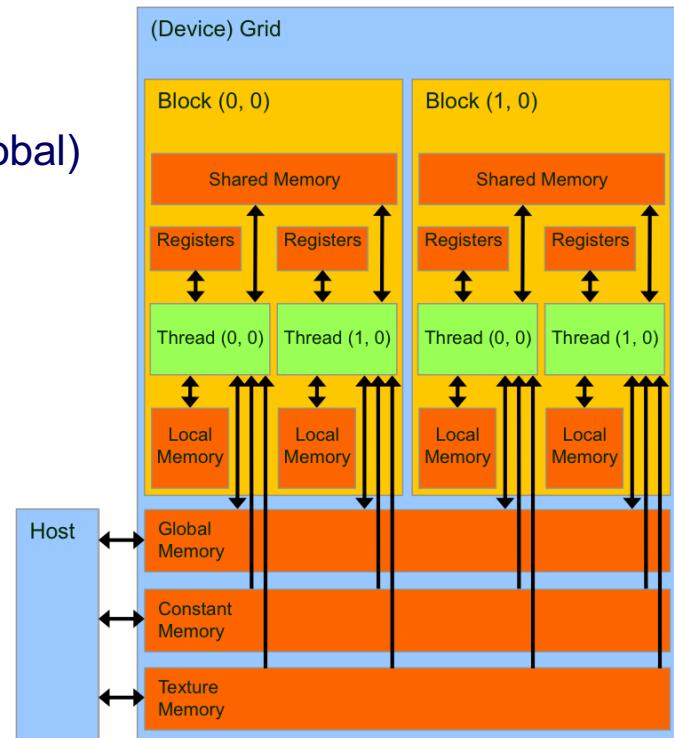
# A more complete view of the device memory

## ■ Each thread can:

- R/W per-thread registers and local memory
- R/W per-block shared memory (100 times faster than global)
- R/W per-grid global memory
- Read only per-grid constant and texture memories

## ■ Host code can:

- R/W global, constant, and texture memories



# A more complete view of the device memory

## ■ Shared memory

- Within a block, used to share data among threads
- Allocated per block, data is not visible to threads in other blocks

## ■ Declared using `__shared__`

```
#define N 1000

__global__ kernel() {
    __shared__ int A[N]; // declared inside kernel bodies.

    // Scope: block. Lifetime: kernel call
    int tid = threadIdx.x;
    A[tid]++;
    ...
}
```

# CUDA Basics

## ■ Outline

- GPU architecture
- Devices, kernel definition and offloading
- Blocks, threads and indexing
- Accessing (global) memory
- Cooperating threads and shared memory
- Reductions in CUDA

# Parallel Reduction

## ■ Common and important data parallel primitive

- Easy to implement in CUDA
  - ✓ Harder to get it right

## ■ Serves as a great optimization example

- We'll walk step by step through 7 different versions
- Demonstrates several important optimization strategies

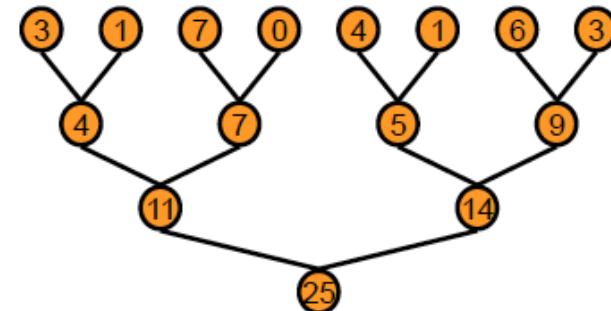
# Parallel Reduction

- Tree-based approach used within each thread block

- Need to be able to use multiple thread blocks

- To process very large arrays
- To keep all multiprocessors on the GPU busy
- Each thread block reduces a portion of the array

- But how do we communicate partial results between thread blocks?



# Problem: Global Synchronization

- If we could synchronize across all thread blocks, could easily reduce very large arrays, right?
  - Global sync after each block produces its result
  - Once all blocks reach sync, continue recursively
- But CUDA has no global synchronization. Why?
  - Expensive to build in hardware for GPUs with high processor count
  - Would force programmer to run fewer blocks (no more than # multiprocessors \* # resident blocks / multiprocessor) to avoid deadlock, which may reduce overall efficiency
- Solution: decompose into multiple kernels
  - Kernel launch serves as a global synchronization point
  - Kernel launch has negligible HW overhead, low SW overhead

# Reduction #1: Interleaved Addressing (I)

```
__global__ void reduce0(int* g_idata, int* g_odata){  
    extern __shared__ int sdata[];  
    // each thread loads one element from global to shared mem  
    unsigned int tid = threadIdx.x;  
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;  
    sdata[tid] = g_idata[i]; Copy to shared memory  
    __syncthreads();  
    // do reduction in shared mem  
    for(unsigned int s=1; s < blockDim.x; s *= 2) {  
        if(tid % (2*s) == 0){ Reduction in shared memory  
            sdata[tid] += sdata[tid + s];  
        }  
        __syncthreads(); Thread synchronization  
    }  
    // write result for this block to global mem  
    if(tid == 0) g_odata[blockIdx.x] = sdata[0]; Copy back to global memory  
}
```

# Reduction #1: Interleaved Addressing (II)

```
// PRE:  
// dA is an array allocated on the GPU  
// N <= len(dA) is a power of two (N >= BLOCKSIZE)  
// POST: the sum of the first N elements of dA is returned  
template<size_t blockSize, typename T>  
T CPUReduction(T* dA, size_t N)  
{  
    T tot = 0.;  
    size_t n = N;  
    size_t blocksPerGrid = std::ceil((1.*n) / blockSize);  
  
    T* g_out1, g_out2;  
    cudaMalloc(&g_out1, sizeof(T) * blocksPerGrid);  
    cudaMalloc(&g_out2, sizeof(T) * blocksPerGrid);  
  
    T* from = dA;  
    T* to = g_out1;  
  
    // Code continues on next slide !!!
```

# Reduction #1: Interleaved Addressing (III)

```
do
{
    blocksPerGrid    = std::ceil((1.*n) / blockSize);
    reduce0<blockSize><<<blocksPerGrid, blockSize>>>(from, to);
    tmp=from; from = to;
    to = tmp;
    n = blocksPerGrid;
} while (n > blockSize);

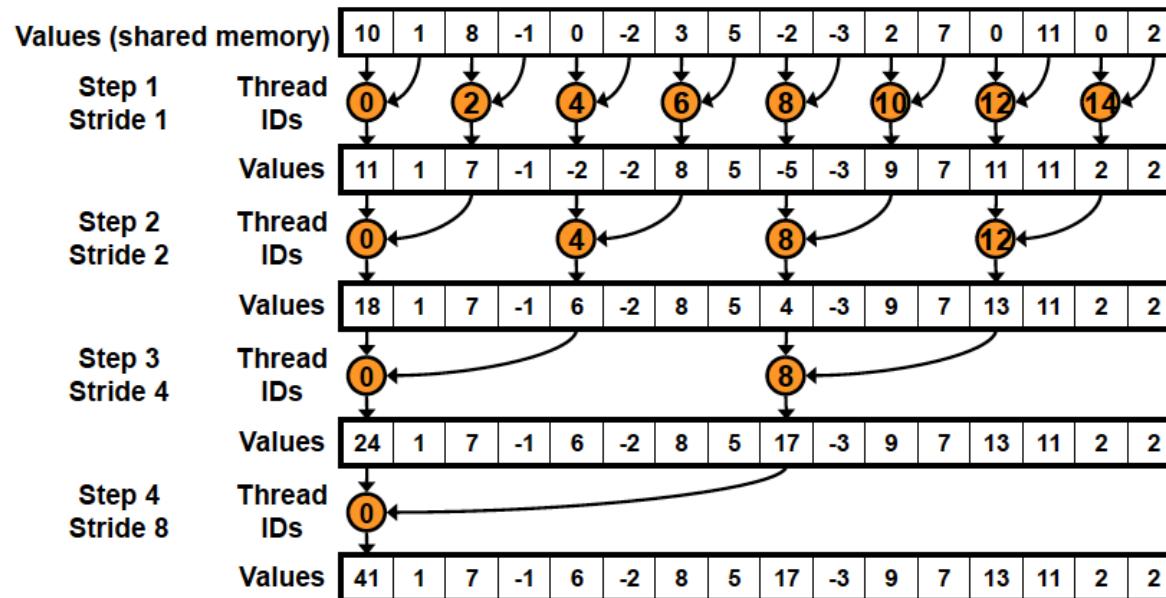
if (n > 1)
    reduce0<blockSize><<<1, blockSize>>>(tmp, tmp);

cudaDeviceSynchronize();

cudaMemcpy(&tot, tmp, sizeof(T), cudaMemcpyDeviceToHost);

cudaFree(tmp);
return tot;
}
```

# Parallel Reduction: Interleaved Addressing



# Reduction #1: Interleaved Addressing

```
__global__ void reduce0(int* g_idata, int* g_odata){  
    extern __shared__ int sdata[];  
    // each thread loads one element from global to shared mem  
    unsigned int tid = threadIdx.x;  
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;  
    sdata[tid] = g_idata[i];  
    __syncthreads();  
    // do reduction in shared mem  
    for(unsigned ints=1; s < blockDim.x; s *= 2) {  
        if (tid % (2*s) == 0){  
            sdata[tid] += sdata[tid + s];  
        }  
        __syncthreads();  
    }  
    // write result for this block to global mem  
    if(tid == 0) g_odata[blockIdx.x] = sdata[0];  
}
```

Problem: highly divergent warps are very inefficient, and % operator is very slow

## Reduction #2: Interleaved Addressing

- Just replace divergent branch in inner loop:

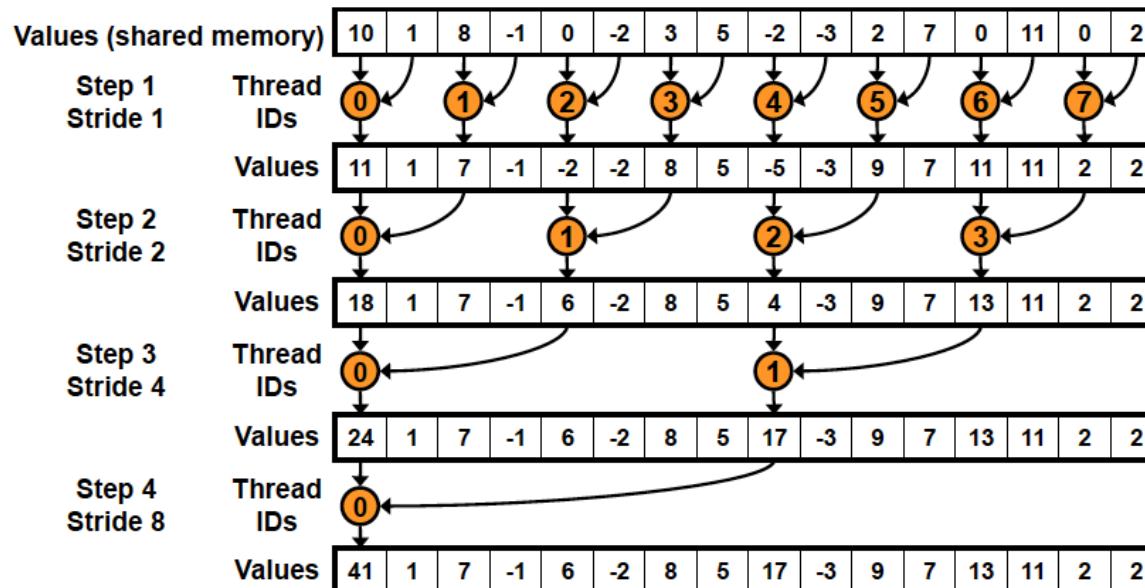
```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    if (tid % (2*s) == 0) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

- With strided index and non-divergent branch:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

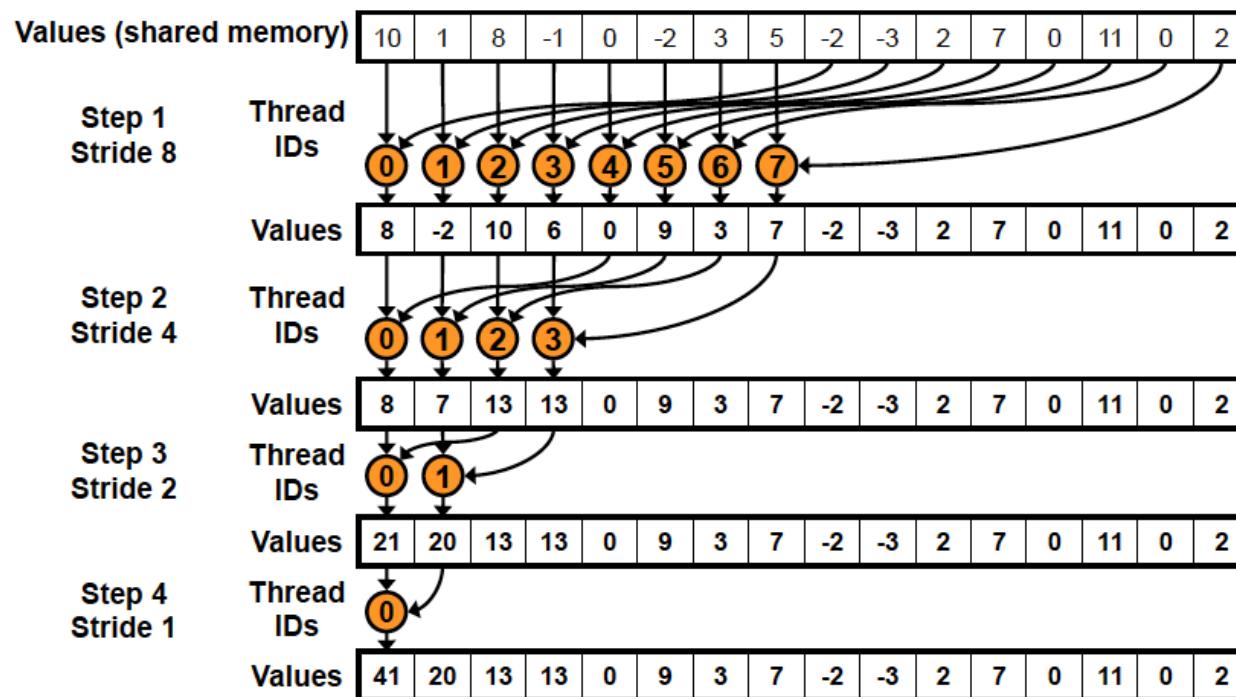
# Parallel Reduction: Interleaved Addressing

## ■ New Problem: Shared Memory Bank Conflicts



# Parallel Reduction: Sequential Addressing

## ■ Sequential Addressing is conflict free



# Reduction #3: Sequential Addressing

- Just replace strided indexing in inner loop:

```
for(unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

- With reversed loop and threadID-based indexing:

```
for(unsigned int s = blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

# Idle Threads

## ■ Problem: Half of the threads are idle on first loop iteration!

- This is wasteful...

```
for(unsigned int s= blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

# Reduction #4: First Add During Load

- Halve the number of blocks, and replace single load:

```
// each thread loads one element from global to shared mem
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*blockDim.x+ threadIdx.x;
sdata[tid] = g_idata[i];
__syncthreads();
```

- With two loads and first add of the reduction:

```
// perform first level of reduction,
// reading from global memory, writing to shared memory
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

# Unrolling the Last Warp

- As reduction proceeds, # “active” threads decreases
  - When  $s \leq 32$ , we have only one warp left
- Instructions are SIMD synchronous within a warp
- That means when  $s \leq 32$ :
  - We don’t need to `__syncthreads()`
  - We don’t need “if ( $tid < s$ )” because it doesn’t save any work
- Let’s unroll the last 6 iterations of the inner loop

# Reduction #5: Unroll the Last Warp

## ■ Unroll latest 6 iterations

```
__device__ void warpReduce(volatile int* sdata, int tid) {  
    sdata[tid] += sdata[tid + 32];  
    sdata[tid] += sdata[tid + 16];  
    sdata[tid] += sdata[tid + 8];  
    sdata[tid] += sdata[tid + 4];  
    sdata[tid] += sdata[tid + 2];  
    sdata[tid] += sdata[tid + 1];  
}
```

Note: This saves useless work in all warps, not just the last one!

Without unrolling, all warps execute every iteration of the for loop and if statement

## ■ Then

```
for(unsigned int s=blockDim.x/2; s>32; s>>=1) {  
    if (tid < s)  
        sdata[tid] += sdata[tid + s];  
    __syncthreads();  
}  
if (tid < 32) warpReduce(sdata, tid);
```

# Reduction #6: Template the unrolling with block size

## ■ Specify block size as a function template parameter:

- Function reduce5 code:

```
if (blockSize >= 512)
    if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads(); }
if (blockSize >= 256)
    if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads(); }
if (blockSize >= 128)
    if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads(); }
if (tid < 32) warpReduce<blockSize>(sdata, tid);
```

```
template <unsigned int blockSize>
__global__ void reduce5(int *g_idata, int *g_odata)
```

## ■ Warp-level function:

```
template <unsigned int blockSize>
__device__ void warpReduce(volatileint* sdata,int tid) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
}
```

Note: all code in RED will be evaluated at compile time.  
Results in a very efficient inner loop!

# Invoking Template Kernels

## ■ Don't we still need block size at compile time?

- Nope, just a switch statement for 10 possible block sizes:

```
switch (threads)
{
    case 512:
        reduce5<512><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 256:
        reduce5<256><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 128:
        reduce5<128><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 64:
        reduce5< 64><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 32:
        reduce5< 32><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 16:
        reduce5< 16><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 8:
        reduce5< 8><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 4:
        reduce5< 4><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 2:
        reduce5< 2><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 1:
        reduce5< 1><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
}
```

# Reduction #7: Multiple Adds / Thread

## ■ Replace load and add of two elements:

```
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

## ■ With a while loop to add as many as necessary:

```
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockSize*2) + threadIdx.x;
unsigned int gridSize = blockSize*2*gridDim.x;
sdata[tid] = 0;
while (i < n) {
    sdata[tid] += g_idata[i] + g_idata[i+blockSize];
    i += gridSize;      Note: gridSize loop
}                                stride to maintain
__syncthreads();                  coalescing
```

# Putting it all together

```
template <unsigned int blockSize>
__device__ void warpReduce(volatile int *sdata, unsigned int tid) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
}

template <unsigned int blockSize>
__global__ void reduce6(int *g_idata, int *g_odata, unsigned int n) {
    extern __shared__ int sdata[];
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*(blockSize*2) + tid;
    unsigned int gridSize = blockSize*2*gridDim.x;
    sdata[tid] = 0;
    while (i < n){sdata[tid] += g_idata[i] + g_idata[i+blockSize]; i += gridSize; }
    __syncthreads();
    if (blockSize >= 512) { if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads(); }
    if (blockSize >= 256) { if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads(); }
    if (blockSize >= 128) { if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads(); }
    if (tid < 32) warpReduce(sdata, tid);
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

# CUDA BASICS