

SURNAME, NAME: ..... DNI/NIF:.....

The duration of the exam is **75 minutes**.The **grades of the exam** will be published on the **12<sup>th</sup> June 2020****Exercise 1 (0,5 point)**

a) Many parallel languages have a barrier construct, e.g. MPI\_Barrier in MPI or #pragma omp barrier in OpenMP. Define barrier.

b) How does one implement a global barrier in CUDA? Explain your answers. **Note:** This is a trick question. There are no barrier constructs in the programming language. The question asks, when do all parallel execution threads synchronize in CUDA?

**Exercise 2 (1 point)**

Given the following C code with OpenMP pragmas:

```
#include #include
#define NUM_OF_COLUMNS 6
#define NUM_OF_ROWS (3*(NUM_OF_COLUMNS - 1))

int whichThread[NUM_OF_ROWS][NUM_OF_COLUMNS];

void fillColumn(int j) {
    int i;
    #pragma omp for
    for (i = 0; i < NUM_OF_ROWS; i++)
        whichThread[i][j] = omp_get_thread_num(); }

int main() {
    int i, j;
    for (i = 0; i < NUM_OF_ROWS; i++) // initialize the array
        for (j= 0; j < NUM_OF_COLUMNS; j++) whichThread[i][j] = -1;

    #pragma omp parallel num_threads(NUM_OF_COLUMNS - 1)
        fillColumn(0);

    #pragma omp parallel for num_threads(NUM_OF_COLUMNS - 1)
        for (j = 1; j < NUM_OF_COLUMNS; j++) fillColumn(j);

    for (i = 0; i < NUM_OF_ROWS; i++) // print out the results
        for (j= 0; j < NUM_OF_COLUMNS; j++) printf(" %2d ", whichThread[i][j]);

    printf("\n"); } return 0;
}
```

SURNAME, NAME: ..... DNI/NIF:.....

a) What is the output of this program if it is compiled **without** the -fopenmp compiler flag? Briefly explain why.

b) What is the output of this program if it is compiled **with** the -fopenmp compiler flag? Explain, very carefully, why the output differs from the output of the serial version of the program.

**Exercise 3 (1,5 points)**

a) Parallelize the following code using OpenMP pragmas. Be sure to explicitly specify the "schedule" options that should be used for better performance.

```
for (i=1; i<N; i++) {
    for (j=1; j<i; j++) {
        C[i] *= A[i][j] + B[i][j];
    }
}
```

b) Let be  $T_c$  the cost of executing one iteration of the j-loop, Let be  $T$  the number of threads to be used. Let be  $N$  the number of iterations in the i-loop. Let be  $Fi(th)$  and  $Li(th)$  the first and last iteration assigned to thread  $th$ . Let be  $Ni(th)$  the total number of iterations assigned to thread  $th$ . Let be  $W(th)$  the total amount of work performed by thread  $th$  if the i-loop were executed in parallel under a **STATIC** scheduling. Assuming  $T$  divides  $N$ , give an expression to model all the previous defined variables:

 $Ni(th) =$  $Fi(th) =$  $Li(th) =$  $W(th) =$

SURNAME, NAME: ..... DNI/NIF:.....

**Exercise 4 (0,5 point)**

Fill the table with any of the two options (device or host) in each case for the function prototype:

Keyword	Executed on the:	Only callable from the:
<code>__device__ void Function()</code>		
<code>__global__ void Function()</code>		
<code>__host__ void Function()</code>		

**Exercise 5 (0,5 points)**

Given a thread organization in the form of a **2D grid of 1D blocks** of threads, use the CUDA built-in variables (gridDim, blockDim, blockIdx, threadIdx) to compute the global thread ID:

**Exercise 6 (1,5 points)**

The following code snippet corresponds to a code skeleton for reduction operations in CUDA. Assume *SharedData* is a vector allocated in shared memory and stores the data for the reduction.

```
for (unsigned int j=blockDim.x >> 1; j>0; j>>=1)
{
    if (tid < j)
        SharedData[tid] += SharedData[tid+j];
    __syncthreads();
}
```

One possible optimization is making a specialized code version for an specific value of blockDim.x = 256:

```
if (tid < 128) SharedData[tid] += SharedData[tid+128]; __syncthreads();
if (tid < 64) SharedData[tid] += SharedData[tid+64]; __syncthreads();
if (tid < 32) {
    SharedData[tid] += SharedData[tid+32];
    SharedData[tid] += SharedData[tid+16];
    SharedData[tid] += SharedData[tid+8];
    SharedData[tid] += SharedData[tid+4];
    SharedData[tid] += SharedData[tid+2];
    SharedData[tid] += SharedData[tid+1];
}
```

This loop unrolling eliminates the call to `__syncthreads()` when there are fewer than 32 active threads in a thread block.

a) Why is `__syncthreads()` necessary in general, e.g. in the loop prior to unrolling? What purpose does it serve?

SURNAME, NAME: ..... DNI/NIF:.....

b) Why is it safe to eliminate `__syncthreads()` in the unrolled loop for fewer than 32 threads? What CUDA principle does this demonstrate?

**Exercise 7 (2,5 point)**

Sketch a CUDA program for adding the elements of a two vectors, A and B and store the result in a vector C, of size N. Your code should include all key CUDA API calls and the usage of shared memory, if necessary.

```
__global__ void Addition(int *a_d, int *b_d, int *c_d, int N)
{
    /* GPU code */

}

int main(void)
{
    const int n = "very large number";
    int a[n], b[n], c[n];

    int *a_dev, *b_dev, *c_dev;

    /* CPU code */

}
```

SURNAME, NAME: ..... DNI/NIF:.....

**Exercise 8 (1,5 points)**

Sketch an MPI program for adding the elements of a vector, V, of size N. Your code should include all key MPI API calls. Assume P processors. Assume the content of the V vector initially is stored in MPI process 0.

```
int main (int argc, char *argv[])
{
    int *a, *b, *c;
    int total_proc; // total nuber of processes
    int rank; // rank of each process
    long long int n_per_proc; // elements per process
    long long int i, n;
    unsigned int MASTER=0;

    MPI_Status status;

    // Initialization of MPI environment
    MPI_Init (&argc, &argv);
    MPI_Comm_size (MPI_COMM_WORLD, &total_proc);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);

    int *ap, *bp, int *cp;

    if (rank == MASTER) {
        a = (int *) malloc(sizeof(int)*n);
        b = (int *) malloc(sizeof(int)*n);
        c = (int *) malloc(sizeof(int)*n);

        MPI_Bcast (&n, _____);
        n_per_proc = n/total_proc;
        MPI_Bcast (_____, _____);

        MPI_Scatter(a, _____);
        MPI_Scatter(_____, _____);

        for(i=0;i<n_per_proc;i++) cp[i] = ap[i]+bp[i];

        MPI_Gather(_____, _____);
    } else { // Non-master tasks
        MPI_Bcast(_____, _____);
        MPI_Bcast(_____, _____);

        ap = _____;
        bp = _____;
        cp = _____;

        MPI_Scatter(_____, _____);
        MPI_Scatter(_____, _____);

        for(i=0;i<n_per_proc;i++) cp[i] = ap[i]+bp[i];

        MPI_Gather(_____, _____);
    }

    MPI_Finalize();
    return 0;
}
```

SURNAME, NAME: ..... DNI/NIF:.....

**Exercise 9 (0,5 point)**

a) In MPI, what is a process rank?

b) In MPI you set the number of processes when you write the source code. **True** or **False**:

c) Explain if the following MPI code segment is correct or not, and why:

Process 0 executes:

```
MPI_Recv( &yourdata, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD, &status);  
MPI_Send( &mydata, 1, MPI_FLOAT, 1, tag, MPI_COMM_WORLD);
```

Process 1 executes:

```
MPI_Recv( &yourdata, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD, &status);  
MPI_Send( &mydata, 1, MPI_FLOAT, 0, tag, MPI_COMM_WORLD);
```