

# OPENMP BASICS

# OpenMP Basics

## ■ Outline

- Overview
- Execution Model
- Threads and Data Access
- Application Program Interface
- Thread Synchronization
- Memory Consistency

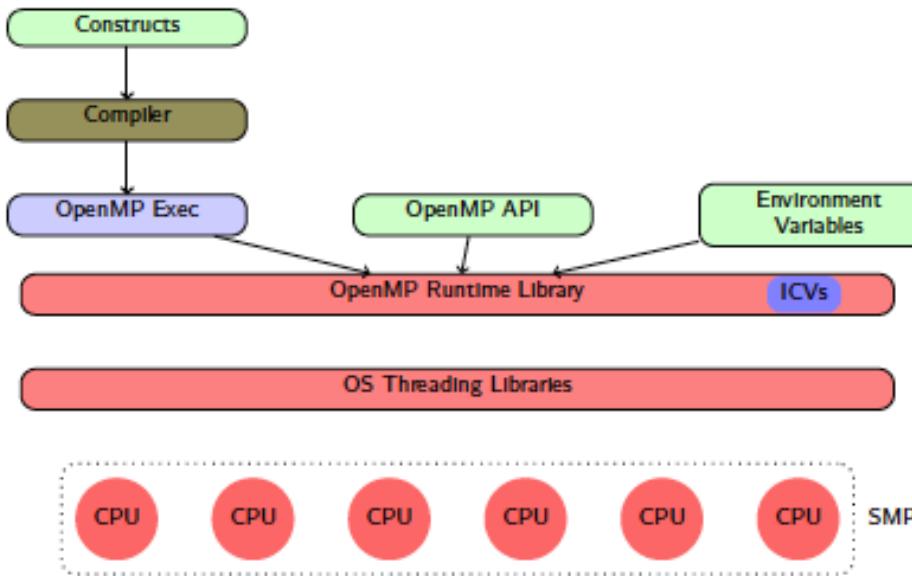
# OpenMP Basics

## ■ What is OpenMP?

- It's an API extension to the C, C++ and Fortran languages to write parallel programs for shared memory machines
- Current version is 5.0 (November 2018)
- Supported by most compiler vendors
  - ✓ Intel, IBM, HP, GCC,...
- Maintained by the Architecture Review Board (ARB), a consortium of industry and academia
- This mini-tutorial just covers part of the specification, for a complete reference please consult the documentation online
  - ✓ <http://www.openmp.org>

# OpenMP Basics

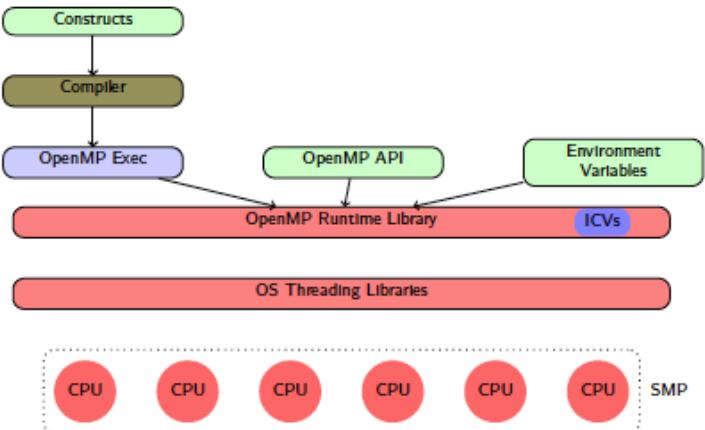
## ■ OpenMP Components:



# OpenMP Basics

## ■ OpenMP Components:

- Constructs
  - ✓ These form the major elements of OpenMP programming
    - Create threads
    - Share the work amongst threads
    - Synchronize threads and memory
- Library routines
  - ✓ To control and query the parallel execution environment (internal control variables - ICVs)
- Environment variables
  - ✓ The execution environment can also be set before the program execution is started



# OpenMP Directives

## ■ OpenMP directives syntax

- In C/C++, through a compiler directive:
  - ✓ `#pragma omp [optional clauses]`
- OpenMP syntax is ignored if specific flag is not set for the compiler

## ■ Most directives apply to:

- A block of one or more statements
  - ✓ One entry point, one exit point
    - No branching in or out allowed

# OpenMP Headers and Macros

## ■ C/C++ only

- `omp.h` contains the API prototypes and data types definitions
- The `_OPENMP` macro is defined by OpenMP enabled compiler
  - ✓ Allows conditional compilation of OpenMP

```
// Conditional code for OpenMP

#ifndef _OPENMP
    // OpenMP code
#else
    // Serial, original code without OpenMP
#endif
```

# Performance Modeling

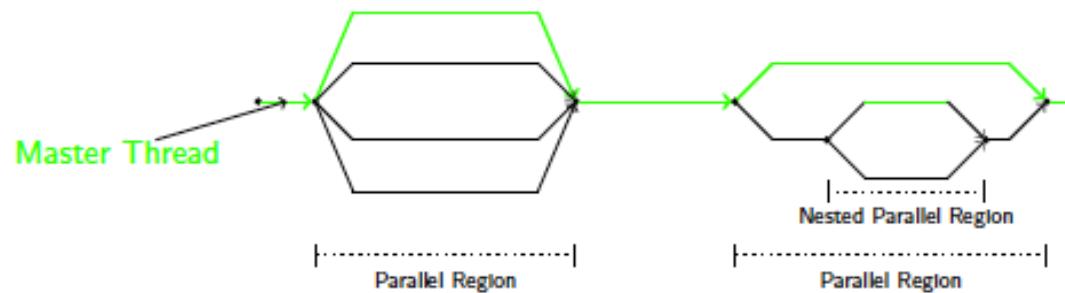
## ■ Outline

- Overview
- Execution Model
- Threads and Data Access
- Application Program Interface
- Thread Synchronization
- Memory Consistency

# OpenMP Execution Model

## ■ Fork-join model

- OpenMP uses a fork-join model
  - ✓ The master thread spawns a team of threads that joins at the end of the parallel region
  - ✓ Threads in the same team can collaborate to do work



# OpenMP Memory Model

## ■ OpenMP defines a relaxed memory model

- Threads can see different values for the same variable
- Memory consistency is only guaranteed at specific points
- Luckily, the default points are usually enough
  - ✓ If not ... there is a mechanism to guarantee it! (described at the end of Part I)
- Variables can be shared or private to each thread

# Performance Modeling

## ■ Outline

- Overview
- Execution Model
- Threads and Data Access
- Application Program Interface
- Thread Synchronization
- Memory Consistency

# Parallelism Definition

## ■ The parallel construct

- Directive

✓ `#pragma omp parallel [ clauses ]  
  – structured block`

- Where some of the clauses are:

✓ `num_threads(expression)`  
✓ `if(expression)`  
✓ `shared(var-list)`      → Coming shortly !  
✓ `private(var-list)`  
✓ `firstprivate(var-list)`  
✓ `reduction(operator:var-list)`      → We will see it later!

# Parallelism Definition

## ■ The **parallel** construct

- Specifying the number of threads
  - ✓ The `nthreads-var` ICV is used to determine the number of threads to be used for encountered parallel regions
  - ✓ It is a list of positive integer values, its first element specifying the number of processors for the next nesting level
  - ✓ When a parallel construct is encountered, and the generating task's `nthreads-var` list contains multiple elements, the generated task(s) inherit the value of `nthreads-var` as the list obtained by deletion of the first element
  - ✓ If the generating task's `nthreads-var` list contains a single element, the generated task(s) inherit that list as the value of `nthreads-var`

# Parallelism Definition

## ■ The **parallel** construct

- Specifying the number of threads
  - ✓ The `nthreads-var` list can be defined from the execution command line by setting the `OMP_NUM_THREADS` environment variable
  - ✓ Example: `setenv OMP_NUM_THREADS 4,3,2`
  - ✓ After that, the `threads-var` list can be modified through:
    - the `omp_set_num_threads` API, which sets the value of the first element of the current list
    - the `num_threads` clause, which causes the implementation to ignore the ICV and use the value of the clause for that region

# Parallelism Definition

## ■ The **if** clause

- Avoiding parallel regions
  - ✓ Sometimes we only want to run in parallel under certain conditions (e.g.: enough input data, not running already in parallel, ...)
- The if clause allows to specify an expression. When evaluates to false the parallel construct will only use 1 thread
- Note that still creates a new team and data environment

# Parallelism Definition

## ■ The **parallel** construct

- How many threads are used in each parallel region below?

```
void main ( ) {  
  
#pragma omp parallel  
.  
.  
.omp_set_num_threads ( 2 );  
  
#pragma omp parallel  
.  
.  
.#pragma omp parallel num_threads (random()%4+1) if ( 0 )  
.  
.  
  
}
```

# Data Sharing Attributes

## ■ The shared clause

- When a variable is marked as shared, the variable inside the construct is the same as the one outside the construct
  - ✓ In a parallel construct this means all threads see the same variable
    - But not necessarily the same value!
- Usually need some kind of synchronization to update them correctly
  - ✓ OpenMP has consistency points at synchronizations
- By default, variables are implicitly shared

# Data Sharing Attributes

- The **private** clause
- When a variable is marked as **private**, the variable inside the construct is a *new* variable of the same type with an *undefined* value
  - In a parallel construct this means all threads have a different variable
  - Can be accessed without any kind of synchronization

# Data Sharing Attributes

## ■ The `firstprivate` clause

- When a variable is marked as `firstprivate`, the variable inside the construct is a `new` variable of the same type but initialized to the original variable value
  - ✓ In a parallel construct this means all threads have a different variable with the same initial value
  - ✓ Can be accessed without any kind of synchronization

# Data Sharing Attributes

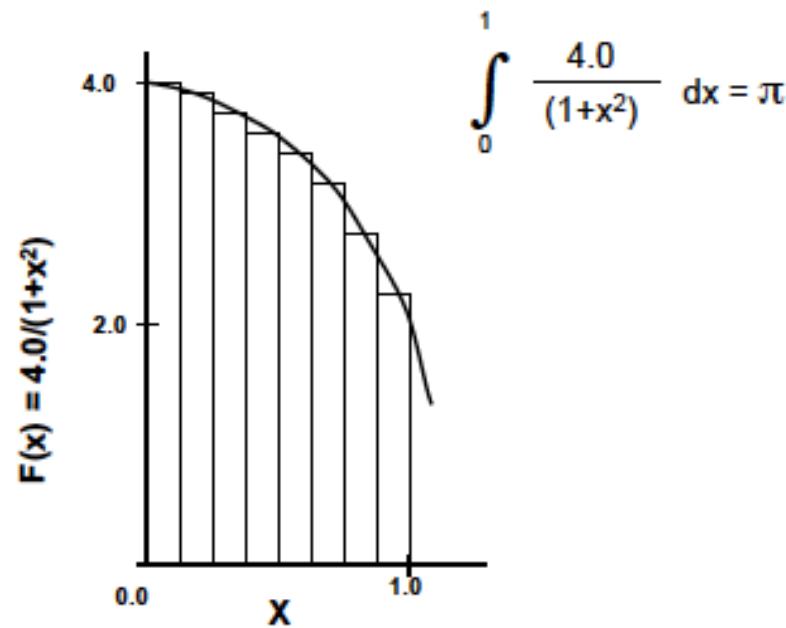
## ■ Exercise

- What does appear on the screen if `xxxxxx` is `shared(x)`, `private(x)` or `firstprivate(x)`?

```
int x=1;
#pragma omp parallel XXXXXX num_threads(2)
{
    x++;
    printf("%d\n", x) ;
}
printf("%d\n", x) ;
```

# Example: Computation of Pi

- Implementation of PI computation with the following numerical recipe:

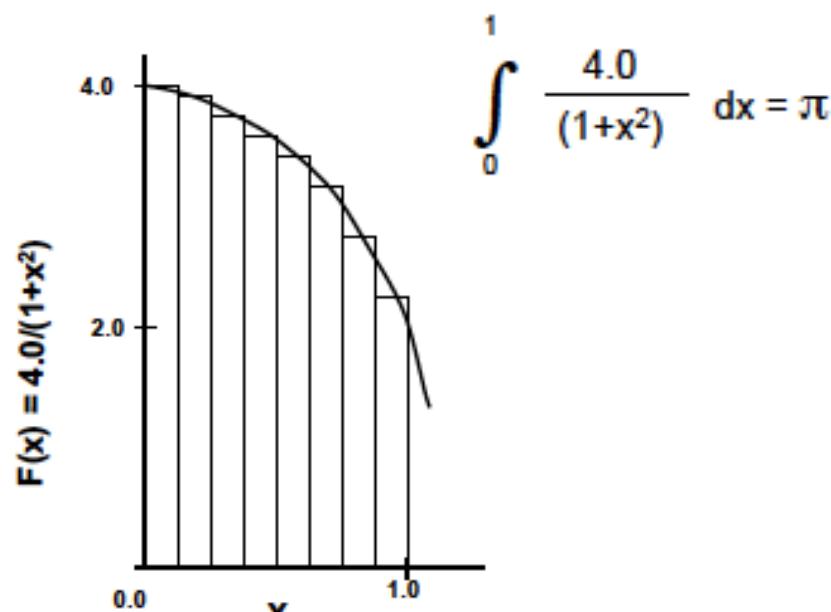


```
static long num_steps = 100000;
double step;

void main ( )
{
    int i;
    double x , pi , sum = 0.0;
    step = 1.0 / (double) num_steps ;
    for (i=1; i<=num_steps; i++) {
        x = (i-0.5)*step;
        sum = sum + 4.0/(1.0+x*x );
    }
    pi = step*sum;
}
```

# Example: Computation of Pi

- Implementation of PI computation with the following numerical recipe:



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

```
static long num_steps = 100000;
double step;
#include <omp.h>
#define NUM_THREADS(2)

void main ( )
{
    int i
    double x , pi , sum = 0.0;
    step = 1.0 / (double) num_steps ;

    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel private(x, i)
    {
        for (i=1; i<=num_steps; i++) {
            x = (i-0.5)*step;
            sum = sum + 4.0/(1.0+x*x );
        }
    }
    pi = step*sum;
}
```

# Performance Modeling

## ■ Outline

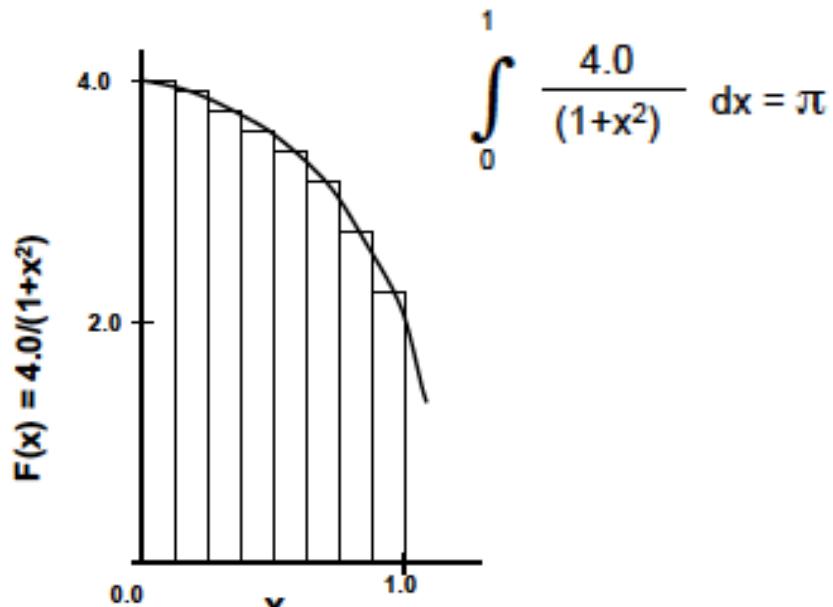
- Overview
- Execution Model
- Threads and Data Access
- Application Program Interface
- Thread Synchronization
- Memory Consistency

# Some API routines

- `int omp_get_num_threads()`
  - Returns the number of threads in the current team. Returns 1 if outside a parallel region
- `int omp_get_thread_num()`
  - Returns the id of the thread in the current team. The id between 0 and `omp_get_num_threads() - 1`
- `void omp_set_num_threads()`
  - Sets the number of threads to be used in parallel regions at the next nesting level
- `int omp_get_max_threads()`
  - Returns the number of threads that could be used in parallel regions at the next nesting level
- `double omp_get_wtime()`
  - Returns the number of seconds since an arbitrary point in the past

# Example: Computation of Pi

## ■ Distribute iterations according to openmp id



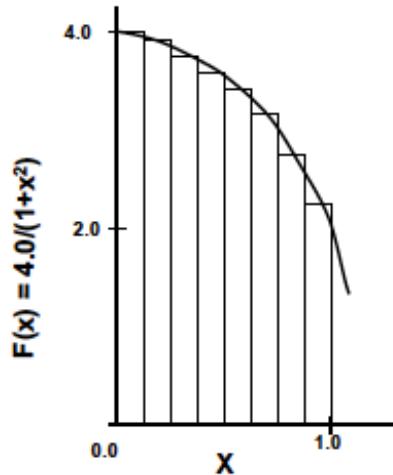
$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

```
static long num_steps = 100000;
double step;
#include <omp.h>
#define NUM_THREADS(2)

void main ( )
{
    int i
    double x , pi , sum = 0.0;
    step = 1.0 / (double) num_steps ;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel private(x, i)
    {
        id = omp_get_thread_num( ) ;
        for (i=id+1; i<=num_steps; i=i+NUM_THREADS) {
            x = (i-0.5)*step;
            sum = sum + 4.0/(1.0+x*x );
        }
        pi = step*sum;
    }
}
```

# Example: Computation of Pi

## ■ Measuring execution time



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

```
static long num_steps = 100000;
double step;
#include <omp.h>
#define NUM_THREADS(2)

void main ( )
{
    int i
    double x , pi , sum = 0.0;
    double TimeStart, TimeStop;

    TimeStart = omp_get_wtime();

    step = 1.0 / (double) num_steps ;

    omp_set_num_threads(NUM_THREADS);

#pragma omp parallel private(x, i)
    {
        id = omp_get_thread_num( );
        for (i=id+1; i<=num_steps; i=i+NUM_THREADS) {
            x = (i-0.5)*step;
            sum = sum + 4.0/(1.0+x*x );
        }
        pi = step*sum;
        TimeEnd = omp_get_wtime();
        printf("Wall clock time = %.20f\n", TimeEnd-TimeStart);
    }
}
```

# Performance Modeling

## ■ Outline

- Overview
- Execution Model
- Threads and Data Access
- Application Program Interface
- Thread Synchronization
- Memory Consistency

# Why synchronization ?

## ■ OpenMP follows a shared memory model

- Threads communicate by sharing variables
- Unintended sharing of data causes race conditions (i.e. the execution outcome may change as the threads are scheduled differently)
- Threads need to synchronize to impose some ordering in their sequence of actions

## ■ Some OpenMP synchronization mechanisms:

- `barrier`
- `critical`
- `atomic`
- Use of locks through API

# Thread Barrier

## ■ **#pragma omp barrier**

- Threads cannot proceed past a barrier point until all threads reach the barrier AND all previously generated work is completed
- Some constructs have an implicit barrier at the end
  - E.g.: the parallel construct

## ■ **Example:**

- Forces all `foo` invokations happen before `bar` invokations
- Implicit barrier at the end of the parallel region

```
#pragma omp parallel
{
    foo ( ) ;

    #pragma omp barrier

    bar ( ) ;
}
```

# Exclusive access: critical construct

## ■ `#pragma omp critical [ (name) ] structured block`

- Provides a region of mutual exclusion where only one thread can be working at any given time
- By default all critical regions are the same
- Multiple mutual exclusion regions by providing them with a name
  - ✓ Only those with the same name synchronize

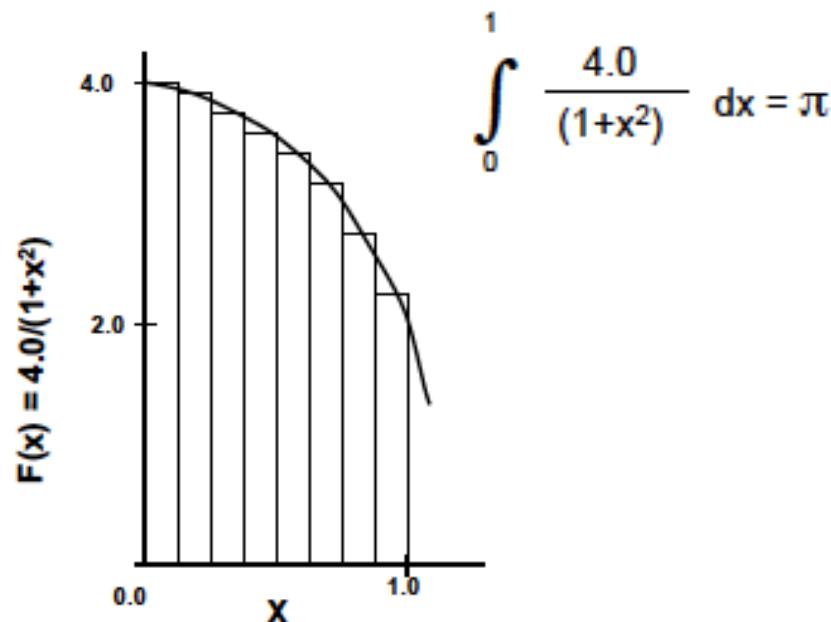
## ■ Example:

- Different names: One thread can update x while another updates y

```
int x=1, y=0;  
#pragma omp parallel num_threads ( 4 )  
{  
#pragma omp critical ( x )  
    x++;  
#pragma omp critical ( y )  
    y++;  
}
```

# Example: Computation of Pi

## ■ `omp critical` usage for sum update



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

```
static long num_steps = 100000;
double step;
#include <omp.h>
#define NUM_THREADS(2)

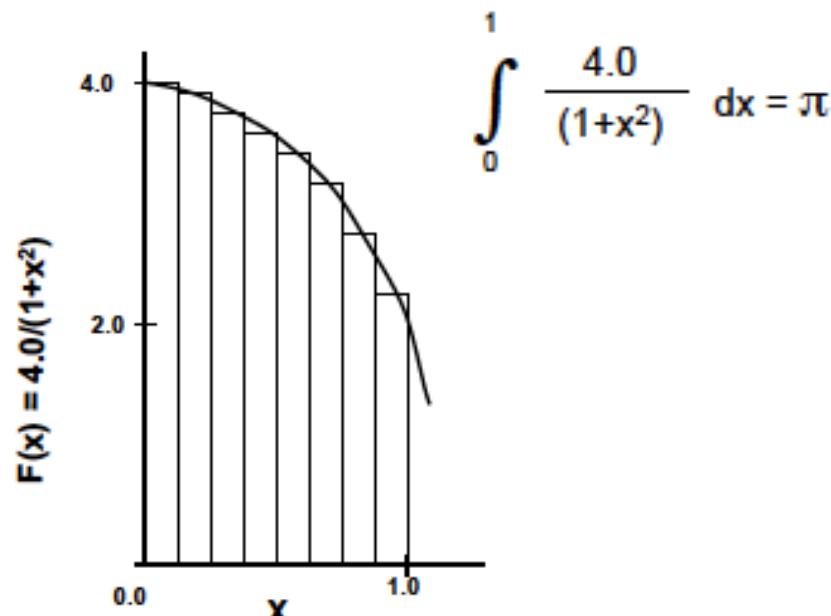
void main ( )
{
    int i
    double x , pi , sum = 0.0;
    step = 1.0 / (double) num_steps ;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel private(x, i)
    {
        id = omp_get_thread_num( ) ;
        for (i=id+1; i<=num_steps; i=i+NUM_THREADS) {
            x = (i-0.5)*step;
#pragma omp critical
            sum = sum + 4.0/(1.0+x*x );
        }
    }
    pi = step*sum;
}
```

# Exclusive access: atomic construct

- `#pragma omp atomic [ update | read | write ] expression`
- Ensures that a specific storage location is accessed atomically, avoiding the possibility of multiple, simultaneous reading and writing threads
  - Atomic updates: `x += 1`, `x = x - foo()`, `x[index[i]]++`
  - Atomic reads: `value = *p`
  - Atomic writes: `*p = value`
- Only protects the read/operation/write
- Usually more efficient than a critical construct
- Other clauses and forms for atomic are allowed in the specification

# Example: Computation of Pi

## ■ `omp atomic` usage for sum update



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

```
static long num_steps = 100000;
double step;
#include <omp.h>
#define NUM_THREADS(2)

void main ( )
{
    int i
    double x , pi , sum = 0.0;
    step = 1.0 / (double) num_steps ;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel private(x, i)
    {
        id = omp_get_thread_num( ) ;
        for (i=id+1; i<=num_steps; i=i+NUM_THREADS) {
            x = (i-0.5)*step;
#pragma omp atomic
            sum = sum + 4.0/(1.0+x*x );
        }
    }
    pi = step*sum;
}
```

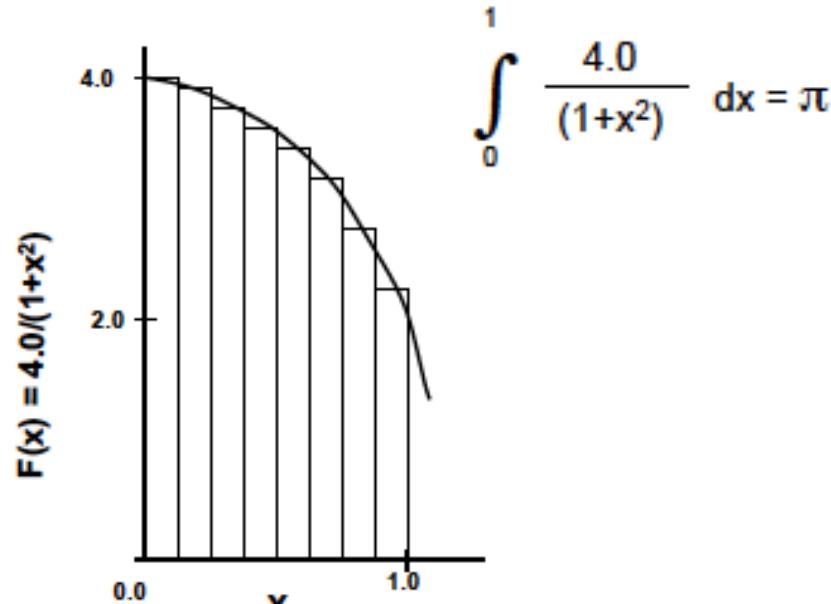
# The reduction clause

- Reduction is a very common pattern where all threads accumulate values into a single variable

- `reduction ( operator : list )`
  - ✓ Valid operators are: +, -, \*, |, ||, &, &&, ^
  - ✓ The compiler creates a private copy of each variable in list that is properly initialized to the identity value
  - ✓ At the end of the region, the compiler ensures that the shared variable is properly (and safely) updated with the partial values of each thread, using the specified operator

# Example: Computation of Pi

## ■ reduction usage for sum update



$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

```
static long num_steps = 100000;
double step;
#include <omp.h>
#define NUM_THREADS(2)

void main ( )
{
    int i
    double x , pi , sum = 0.0;
    step = 1.0 / (double) num_steps ;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel private(x, i) reduction(+:sum)
    {
        id = omp_get_thread_num( ) ;
        for (i=id+1; i<=num_steps; i=i+NUM_THREADS) {
            x = (i-0.5)*step;
            sum = sum + 4.0/(1.0+x*x );
        }
        pi = step*sum;
    }
}
```

# Exclusive access: Locks

## ■ OpenMP provides lock primitives for low-level synchronization

- `omp_init_lock`
  - ✓ Initialize the lock
- `omp_set_lock`
  - ✓ Acquires the lock
- `omp_unset_lock`
  - ✓ Releases the lock
- `omp_test_lock`
  - ✓ Tries to acquire the lock (won't block)
- `omp_destroy_lock`
  - ✓ Frees lock resources

# Exclusive access: Locks

## ■ Example

```
#include <omp.h>
void foo ( )
{
    omp_lock_t lock;
    omp_init_lock(&lock) ;
#pragma omp parallel
{
    omp_set_lock(&lock) ;
    // mutual exclusion region
    omp_unset_lock(&lock);
}
omp_destroy_lock(&lock);
}
```

# Performance Modeling

## ■ Outline

- Overview
- Execution Model
- Threads and Data Access
- Application Program Interface
- Thread Synchronization
- Memory Consistency

# Relaxed consistency memory model

- A thread's temporary view of memory is not required to be consistent with memory at all times
- A value written to a variable can remain in the thread's temporary view until it is forced to memory at a later time
- Likewise, a read from a variable may retrieve the value from the thread's temporary view, unless it is forced to read from memory

# The flush construct

## ■ `#pragma omp flush (list)`

- It enforces consistency between the temporary view and memory for those variables in list
- Synchronization (implicit or explicit) constructs have an associated flush operation

```
#include <omp.h>
void foo ( )
{
    double A[ ];

#pragma omp parallel
{
    // Variable A is not synchronized
    A[ ... ] = ...;

#pragma omp flush(A)

    // Variable A is synchronized
    // All threads see the same values
    ... = A[ ... ];
}
```

# LOOP PARALLELISM IN OPENMP

# Loop Level Parallelism

- Work Sharing Concept
- The Loop Worksharing
- The Single Worksharing

# Work Sharing Concept

- Works-haring constructs divide the execution of a code region among the members of a team
  - Threads **cooperate** to do some work
  - Better way to split work than using thread-ids
  - Lower overhead than using tasks (next section)
    - ✓ But, less flexible
- In OpenMP, there are four work-sharing constructs:
  - loop
  - single
  - sections (not that much used, almost substituted by tasks. We will skip it)
  - workshare (rarely used, we will skip it)

# Loop Work Sharing

## ■ The **for** construct

- `#pragma omp for [ clauses ]  
 ✓ for ( init expr ; test expr ; incr expr )`

## ■ where some possible clauses are:

- `private`
- `firstprivate`
- `reduction`
- `schedule(schedule-kind)`
- `nowait`
- `collapse(n)`

# Loop Work Sharing

## ■ The **for** construct

- The iterations of the loop(s) associated to the construct are divided among the threads of the team.
  - ✓ Loop iterations must be independent
  - ✓ Loops must follow a form that allows to compute the number of iterations
  - ✓ Valid data types for induction variables are: integer types, pointers and random access iterators (in C++)
  - ✓ The induction variable(s) are automatically privatized
  - ✓ The default data-sharing attribute is shared
- It can be merged with the parallel construct:
  - ✓ `#pragma omp parallel for`

# Loop Work Sharing

## ■ Example (Pi computation)

```
#include <omp.h>
static long num_steps = 100000;
double step;
#define NUM_THREADS (2)
void main() {
    int i, id;
    double x, pi, sum;
    step = 1.0 / ( double ) num_steps;
    omp_set_num_threads (NUM_THREADS) ;
#pragma omp parallel for private (x) reduction (+: sum)
    for (i=1; i<=num_steps; i++) {
        x = (i-0.5)*step ;
        sum = sum + 4.0/(1.0+x*x) ;
    }
    pi = sum * step;
}
```

# The schedule clause

- The **schedule** clause determines which iterations are executed by each thread.
  - If no schedule clause is present then is implementation-defined
- There are several possible options as **schedule**:
  - static[,chunk]
  - dynamic[,chunk]
  - guided[,chunk]

# The schedule clause

## ■ static

- The iteration space is broken in chunks of approximately size  $N/\text{num\_threads}$ . Then these chunks are assigned to the threads in a Round-Robin fashion.

## ■ static, N (also called interleaved)

- The iteration space is broken in chunks of size N. Then these chunks are assigned to the threads in a Round-Robin fashion.

## ■ Characteristics of static schedules

- Low overhead
- Good locality (usually)
- Can have load imbalance problems



**Exercise:** check on [openmp.org](http://openmp.org)  
the exact arithmetic to compute  
the iteration distribution!!!

# The schedule clause

## ■ **dynamic, N**

- Threads dynamically grab chunks of N iterations until all iterations have been executed. If no chunk is specified, N = 1.

## ■ **guided, N**

- Variant of dynamic. The size of the chunks decreases as the threads grab iterations, but it is at least of size N. If no chunk is specified, N = 1.

## ■ **Characteristics of dynamic schedules**

- Higher overhead
- Not very good locality (usually)
- Can solve imbalance problems

**Exercise:** check on [openmp.org](http://openmp.org) the description of how iterations are distributed!!!

# The nowait clause

- When a worksharing has a nowait clause then the implicit barrier at the end of the loop is removed.
- This allows to overlap the execution of non-dependent loops/tasks/worksharings

```
#pragma omp for nowait
for (i = 0 ; i < n ; i++)
    v[i] = 0 ;

#pragma omp for
for (i = 0 ; i < n ; i++)
    a[i] = 0 ;
```

# The nowait clause

- Useful to overlap the execution of two (or more) consecutive loops if they have the same static schedule and all have the same number of iterations.

```
#pragma omp for schedule (static, 2) nowait
    for (i = 0 ; i < n ; i++)
        v[i] = ...;

#pragma omp for schedule (static, 2)
    for (i = 0 ; i < n ; i++)
        a[i] = v[i] * v[i] ;
```

# The collapse clause

## ■ Allows to distribute work from a set of n nested loops.

- Loops must be perfectly nested
- The nest must traverse a rectangular iteration space

```
#pragma omp for collapse(2)
for (i = 0 ; i < N; i++)
    for (j = 0 ; j < M; j++)
        foo(i, j);
```

i-loop and j-loop are folded  
and iterations distributed  
among all threads.  
Both i and j are privatized

# The single construct

## ■ Assigns work to a single thread

```
#pragma omp single [ clauses ]  
    structured block
```

- Clauses can be:
  - ✓ `private`
  - ✓ `firstprivate`
  - ✓ `nowait`
- Only one thread of the team executes the structured block
- There is an implicit barrier at the end



# **TASK PARALLELISM IN OPENMP**

# OpenMP Tasking Model

## ■ OpenMP Tasks

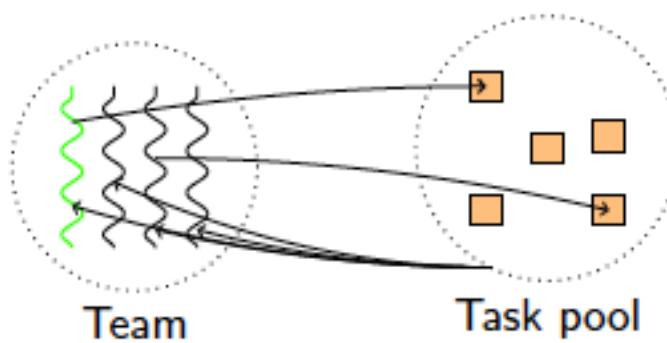
## ■ Task Synchronization

# Task parallelism model

- Tasks are work units whose execution may be deferred
  - In some cases, they can also be executed immediately

- Threads of the team cooperate to execute them
- We distinguish between two essential events

- Task Creation
- Task Execution



# Task Creation: Implicit and Explicit

## ■ Creating tasks

- Implicit
  - ✓ Parallel regions create tasks, as many as threads execution in the team
  - ✓ One implicit task is created and assigned to each thread
- Explicit
  - ✓ Each thread that encounters a task construct
    - Packages the code and data
    - Creates a new explicit task

# Creating (explicit) tasks

## ■ The task construct

```
#pragma omp task [ clauses ]  
    structured block
```

## ■ Where some possible clauses are:

- shared
- private
- firstprivate
  - ✓ Values are captured at creation time
- if(expression)
- final(expression)
- mergeable

# Example: List Traversal

## ■ Code scheme

- We need threads to execute the tasks ...
- ... but not that many! This will generate multiple traversals

```
List L;  
  
L = ...;  
  
...  
#pragma omp parallel  
traverse_list(L);  
  
...  
  
void traverse_list (List* L)  
{  
    Element E;  
    for ( E = L->first; E ; E = E->>next) {  
#pragma omp task  
        process(E);  
    }  
}
```

# Example: List Traversal

## ■ Code scheme

- Using single

✓ One thread creates the tasks of the traversal. The rest (and this one, once task generation is finished) cooperate to execute them

```
...
#pragma omp parallel
{
#pragma omp single
    traverse_list(L);
}
...

void traverse_list (List* L)
{
    Element E;
    for ( E = L->first; E ; E = E->>next) {
#pragma omp task
        process(E);
    }
}
```

# Default task data-sharing attributes

## ■ When no data clauses are specified, some rules apply:

- Global variables are shared
- Variables declared in the scope of a task are private
- The rest are `firstprivate` except when a `shared` attribute can be lexically inherited

# Default task data-sharing attributes

## ■ Exercise

- Global variables are shared
- Variables declared in the scope of a task are private
- The rest are `firstprivate` except when a shared attribute can be lexically inherited

## ■ Explain the scope of the variables

```
int a;
void foo( ) {
    int b, c;
#pragma omp parallel shared (b)
#pragma omp parallel private (b)
{
    int d;
#pragma omp task
    {
        int e;
        a = // shared
        b = // firstprivate
        c = // shared
        d = // firstprivate
        e = // private
    }
}}
```

# Immediate task execution

## ■ The **if** clause

- If the expression of an if clause evaluates to **false**
  - ✓ The encountering task is suspended
  - ✓ The new task is **executed immediately**
    - with its own data environment
    - different task with respect to synchronization
  - ✓ The parent task resumes when the task finishes
  - ✓ Allows implementations to **optimize** task creation

# Immediate task execution (nested)

## ■ The final and mergeable clauses

- **final(expression)**
  - ✓ If the expression of an if clause evaluates to **true**
    - The generated task and all of its child tasks will be **final**
  - ✓ The execution of a **final** task is sequentially included in the generating task (executed immediately)
- **mergeable**
  - ✓ When a **mergeable** clause is present on a task construct, and the generated task is an **included** task, the implementation may generate a **merged task** instead (i.e. no task and context creation for it).

# Immediate task execution (nested)

- final and mergeable tasks (data race!)

```
int fibonacci(int n) {  
  
    int i, j;  
  
    if (n<2) return n;  
  
#pragma omp task shared(i) final(n <= THOLD) mergeable  
    i = fibonacci(n-1);  
  
#pragma omp task shared(j) final(n <= THOLD) mergeable  
    j = fibonacci (n-2);  
  
    ...  
    return (i+j);  
}
```

# OpenMP Tasking Model

## ■ OpenMP Tasks

## ■ Task Synchronization

# Task Synchronization

## ■ There are two types of task barriers:

- **taskwait**

- ✓ Suspends the current task waiting on the completion of child tasks of the current task. The taskwait construct is a stand-alone directive

- **taskgroup**

- ✓ Suspends the current task at the end of structured block waiting on completion of child tasks of the current task and their descendent tasks.

# Task Synchronization: taskwait

```
#pragma omp taskwait

#pragma omp task {}          // T1

#pragma omp task            // T2
{
    #pragma omp task {}      // T3
}

#pragma omp task {}          // T4

#pragma omp taskwait
```

Only T1, T2 and T4 are guaranteed to have finished here!!!

# Taskwait for correct Fibonacci parallelization

```
#pragma omp taskwait
```

```
int fibonacci(int n) {  
  
    int i, j;  
  
    if (n<2) return n;  
  
#pragma omp task shared(i) final(n <= THOLD) mergeable  
    i = fibonacci(n-1);  
  
#pragma omp task shared(j) final(n <= THOLD) mergeable  
    j = fibonacci (n-2);  
  
#pragma omp taskwait  
  
    return (i+j);  
}
```

# Task Synchronization: taskgroup

```
#pragma omp taskgroup
    structured block

#pragma omp task {}          // T1
#pragma omp taskgroup
{
    #pragma omp task           // T2
    {
        #pragma omp task {}    // T3
    }

    #pragma omp task {}        // T4
}

} // taskgroup      Only T2, T3 and T4 are guaranteed to have finished here!!!
```

# Data sharing inside tasks

- In addition one can use `critical` and `atomic` to synchronize the access to shared data inside the task

```
void process ( Element e )
{
    . . .
#pragma omp atomic
    solutions _found++;
    . . .
}
```

# Task dependences

- **Definition of dependences between sibling tasks (i.e. from the same father)**

```
#pragma omp task  [ depend ( in : varlist) ]  
                  [ depend ( out : varlist) ]  
                  [ depend ( inout : varlist) ]
```

- **Task dependences are derived from the dependence type (in, out or inout) and its items in varlist. This list may include array sections**

# Task dependences

- The in dependence-type: the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an out or inout dependence-type list
- The out and inout dependence-types: the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an in, out, or inout dependence-type list.

# Task dependences

## ■ Example: wave-front execution

```
#pragma omp parallel private(i ,j)
#pragma omp single
{
    for (i=1; i<n i++) {
        for (j=1; j<n ; j++) {
            #pragma omp task depend(in : block[i-1][j], block[i][j-1])
                           depend(out : block[i][j])
            foo(i, j);
        }
    }
}
```

# Reduction operations and tasks

## ■ Exercise:

- Check current specification in openmp.org